



**University of
East London**

Pioneering Futures Since 1898

GridGuardian: A Physics-Informed Machine Learning Framework for UK Power Grid Stability Prediction and Explainable Alerting

Fatema Doctor

Student ID: 2604383

A dissertation submitted in partial fulfilment of the requirements for the degree
of

BSc (Hons) in Data Science and Artificial Intelligence

Supervisor: Ms. Dhara Parekh

Module: CN6000

Department of Engineering and Computing
School of Architecture, Computing and Engineering

University of East London

2026

Abstract

Declining system inertia caused by the displacement of synchronous generation with inverter-based renewables has fundamentally altered UK power grid dynamics, reducing the time available for corrective action from minutes to seconds. This dissertation presents **GridGuardian**, a physics-informed machine learning framework designed to provide predictive, explainable early warning of power grid instability. The system employs LightGBM quantile regression to predict 10th–90th percentile frequency bounds 10 seconds in advance, delivering probabilistic risk assessments before critical thresholds are breached.

Key innovations include:

- (1) physics-informed feature engineering incorporating Rate of Change of Frequency (RoCoF), Optimised Swinging Door Algorithm (OpSDA) wind ramp rates, and renewable penetration ratios;
- (2) probabilistic forecasting with uncertainty quantification through quantile regression;
- (3) real-time SHAP-based explanations providing transparent risk drivers for grid operators.

Validation against the August 9, 2019 UK blackout demonstrates the system’s effectiveness: the quantile lower-bound path provided 69 seconds of advance warning before the system frequency reached its catastrophic nadir of 48.8 Hz. While the reactive threshold trigger fired 5 seconds after the initial 49.8 Hz safety threshold was breached, the physics-informed multi-signal fusion provided a proactive window for automated containment systems to act before emergency load-shedding occurred. This engineering limitation regarding reactive triggers was subsequently resolved through the implementation of alert fusion logic (§3.3.1). Model performance metrics include a Pinball Loss of 0.00268 ($\alpha=0.1$ lower bound), Mean Absolute Error (MAE) of 0.021 Hz, and real-time dashboard latency of 0.40 seconds.

The research demonstrates that predictive early warning is achievable for low-inertia grids, with the model showing notable seasonal robustness across summer and winter conditions. GridGuardian represents a proof-of-concept for transitioning grid stability management from purely reactive response toward predictive prevention.

Keywords: power grid stability, low-inertia systems, quantile regression, explainable AI, SHAP, machine learning, renewable energy

Acknowledgements

I would like to express my sincere gratitude to my supervisor, Ms. Dhara Parekh, for her guidance, feedback, and support throughout the research and development of this dissertation.

I am grateful to the University of East London, School of Architecture, Computing and Engineering, for providing the academic environment and resources that made this work possible.

This research would not have been possible without access to publicly available datasets provided by the National Energy System Operator (NESO) via the CKAN open data platform, and meteorological data from the Open-Meteo API. I acknowledge these organisations for their commitment to open data in the energy sector.

Finally, I would like to thank my family and friends for their encouragement throughout this process.

Contents

Abstract	2
Acknowledgements	3
1 Introduction	7
1.1 Background	7
1.2 Problem Statement and Rationale	7
1.3 Research Question and Hypotheses	9
1.4 Research Aim and Objectives	9
1.5 Research Outline	10
2 Literature Review	10
2.1 Introduction	10
2.2 Key Terms and Definitions	10
2.3 Physics of Grid Stability	11
2.3.1 The Swing Equation and System Dynamics	11
2.3.2 The UK Context and August 2019 Blackout	11
2.3.3 Synthetic Inertia and Stability Services	12
2.4 Machine Learning in Power Systems	12
2.4.1 Statistical and Deep Learning Approaches	12
2.4.2 Gradient Boosting and Tree-Based Methods	12
2.4.3 Physics-Informed Neural Networks	12
2.4.4 Hybrid Ensemble Monitoring	13
2.5 Explainable AI in Power Systems	13
2.5.1 SHAP and Feature Attribution	13
2.5.2 Model Interpretability Challenges	13
2.6 Quantile Regression for Probabilistic Forecasting	13
2.6.1 Theoretical Foundations	13
2.6.2 Applications to Power Systems	14
2.7 Gaps in the Literature	14
3 Methodology	15
3.1 Introduction	15
3.2 Research Design	15
3.3 Practical Approach: Implementation and Software Development	16
3.3.1 Software Development Lifecycle	16
3.3.2 Technical Method	16
3.3.3 Technical Method: Hybrid Ensemble Monitoring	17
3.3.4 Multi-Physics Alert Logic	18
3.3.5 Isotonic Quantile Recalibration	18
3.4 Theoretical Approach: Field Research	18
3.4.1 Literature Strategy	18
3.4.2 Quantitative Research Using Secondary Data	18
3.4.3 Data Acquisition and Pre-processing	19
3.5 Challenges and Limitations	19
3.5.1 Inertia Data Granularity	19
3.5.2 RoCoF Acceleration Unit Correction	20
3.5.3 Sensor Micro-Jitter in RoCoF	20
3.5.4 Quantile Calibration Challenges	20
3.6 Ethical and Legal Considerations	21

3.6.1	Data Protection	21
3.6.2	Research Ethics	21
3.6.3	Operational Safety	21
4	Results	22
4.1	Model Performance Metrics	22
4.2	The August 9, 2019 Blackout Reconstruction	23
4.2.1	Multi-Physics Signal Analysis (The Fragility Fingerprint)	24
4.2.2	LSTM Uncertainty Monitor Performance	24
4.3	Feature Importance Analysis	25
4.4	Out-of-Season Validation	27
4.4.1	Analysis of Seasonal Robustness	27
4.4.2	Root Cause Analysis	28
4.4.3	Implications for Generalisability	28
4.5	Dashboard Performance Metrics	29
4.6	Summary	29
5	Analysis and Discussion	30
5.1	Synergistic Architecture: The Hybrid Ensemble Strategy	30
5.1.1	Model Specialization: Forecasting vs. Monitoring	30
5.1.2	Quantile Regression Advantages	31
5.1.3	Interpretability Benefits	32
5.2	Safety-Critical Analysis: The Pessimistic Bias as Feature	32
5.2.1	Calibration Analysis and Safety-Critical Interpretation	32
5.2.2	Fail-Safe Design Philosophy	34
5.2.3	Operational Acceptance	34
5.3	Explainability in Practice: SHAP Analysis of the Blackout	35
5.4	Validation Against Research Objectives	36
5.5	Operational Requirements Assessment	37
5.5.1	Prediction Horizon Sufficiency	37
5.5.2	Comparison to Existing Systems	37
5.5.3	Limitations Preventing Operational Deployment	37
5.6	Theoretical Contributions	39
5.7	Summary	40
6	Conclusion	41
6.1	Summary of Key Findings	41
6.2	Limitations	41
6.2.1	Inertia Data Granularity	42
6.2.2	Single-Season Training	42
6.2.3	Quantile Calibration	42
6.3	Lessons Learned	42
6.3.1	High-Resolution Time-Series Performance	42
6.3.2	Explainability as Operational Requirement	43
6.3.3	Agile Development for Complex Systems	43
6.4	Recommendations for Future Research	43
6.4.1	Full Integration of Half-Hourly Inertia Data	43
6.4.2	Dynamic H in Intervention Simulator	43
6.4.3	Multi-Point Quantile Calibration	43
6.4.4	Cross-Season Training and Multi-Year Validation	44
6.4.5	Formal Ablation Study	44

6.5	Contributions to Knowledge	44
6.6	Concluding Remarks	44
7	References	46
8	Appendices	47
8.1	Appendix A: OpSDA Algorithm Implementation	47
8.2	Appendix B: Full Model Evaluation Tables	49
8.3	Appendix C: Dashboard User Manual and Interface Documentation	51
8.4	Appendix D: Dashboard Architecture & Cache Design	56

1 Introduction

1.1 Background

The United Kingdom’s electricity system is undergoing a profound structural transformation. Following the Climate Change Act 2008 (UK Parliament, 2008) and subsequent net-zero commitments, the nation has systematically retired fossil fuel generation in favour of renewable energy sources. The closure of Ratcliffe-on-Soar on 30 September 2024 marked the end of 142 years of coal-powered electricity in Britain, representing a historic milestone in this transition (National Grid ESO, 2024). Renewables now constitute over 50% of UK electricity generation, with wind power alone contributing approximately 29.5% of total supply (Department for Energy Security and Net Zero, 2024).

However, this decarbonisation success has introduced an unprecedented engineering challenge. Traditional synchronous generators—coal, gas, and nuclear plants—provided inherent physical inertia through their large rotating masses. This stored kinetic energy acted as a natural shock absorber, stabilising grid frequency during generation-demand imbalances (Tielens & Van Hertem, 2016). Inverter-based renewable sources, by contrast, lack this mechanical inertia. As conventional generation retires, system inertia has declined by over 50% since 2010, and National Grid ESO (2023) projects minimum inertia levels falling to 96 GVA · s by 2025 for zero-carbon operation.

This phenomenon, termed the “inertia crisis” by Saleem et al. (2024), fundamentally alters grid dynamics. Lower inertia permits faster rates of change of frequency (RoCoF), reducing the time available for corrective action from minutes to mere seconds. The relationship between inertia and frequency stability is governed by the swing equation, which demonstrates that systems with reduced inertia experience exponentially faster frequency deviations during disturbances (Kundur et al., 1994). Without intervention, small disturbances can cascade into widespread blackouts, threatening the reliability that underpins modern economic activity.

1.2 Problem Statement and Rationale

The economic and societal consequences of grid instability are substantial. National Grid ESO currently expends over £650 million annually on “synthetic inertia” procurement through Stability Pathfinder programmes, contracting grid-scale batteries and other technologies to artificially stabilise the system (Amamra, 2025). These costs represent a significant burden on energy consumers and highlight the economic imperative of improved stability management.

The August 9, 2019 blackout exemplifies the operational risks. At 15:52 UTC, a lightning strike triggered cascading failures at Little Barford gas station and Hornsea One offshore wind farm, resulting in 1,481 MW of instantaneous generation loss (Homan, 2020). Frequency collapsed to 48.8 Hz—exactly the automatic load-shedding threshold—within 10 seconds. The consequences affected 1.1 million customers, paralysed railway networks for days, and caused Ipswich Hospital’s backup generators to fail when they could not synchronise with the collapsing grid (Ofgem & BEIS, 2019).




Figure 1: Grid Stability Phase Portrait: High-resolution trajectory of the August 9 collapse showing the spiral into the critical zone.




Figure 2: August 9, 2019 Frequency Data (Actual Nadir: 48.8 Hz)

Critically, traditional monitoring systems provide only reactive responses. Frequency thresholds trigger only after deviations occur, leaving insufficient time for human operators to implement corrective measures. Automated systems require approximately 1–2 seconds to inject power from Firm Frequency Response (FFR) batteries (Hong et al., 2021). The research aims to develop a predictive capability providing a 10-second lookahead—sufficient time to enable these

automatic responses to activate and mitigate frequency collapse. Validation against the August 2019 event demonstrates that the quantile-based lower bound provides 69 seconds of lead-time before the system reaches catastrophic levels (48.8 Hz nadir), potentially preventing total blackout through timely containment.

The research therefore addresses a timely and significant gap: current grid management lacks predictive, explainable early warning capabilities that could transition stability management from purely reactive threshold-based monitoring toward anticipatory prevention of worst-case outcomes.

Beyond the physical challenge, the implementation of such a system presents a significant software engineering hurdle. Real-time machine learning deployment in safety-critical environments requires sub-second inference latency whilst processing multi-resolution data streams. Standard monolithic application architectures often struggle with “IO-bound” bottlenecks during model loading and feature engineering. This research specifically addresses this through a novel split-cache architecture, separating persistent model resources from reactive temporal data caches to maintain operational readiness in high-velocity grid environments.

1.3 Research Question and Hypotheses

This research investigates the following primary question:

“How can machine learning be employed to provide predictive, explainable early warning of power grid instability in low-inertia, high-renewable energy systems?”

From this question, two competing hypotheses emerge:

Primary Hypothesis (H_1): A physics-informed quantile regression model can predict grid frequency instability 10 seconds in advance with sufficient reliability to serve as an effective early warning system for grid operators.

Null Hypothesis (H_0): Such a model cannot provide reliable predictions beyond existing methods, and machine learning offers no material improvement over traditional threshold-based monitoring.

Validation against the August 9, 2019 blackout data provides a critical empirical test. If the model successfully predicts this documented instability event, H_1 receives substantial support; failure to do so would indicate the need for alternative methodological approaches.

1.4 Research Aim and Objectives

The aim of this research is to design, implement, and validate a physics-informed early warning system for UK power grid instability.

The specific objectives are:

1. **Develop a high-performance data pipeline** utilising the Polars library for real-time processing of 1-second resolution grid frequency and weather data, achieving significant performance improvements over conventional Pandas-based approaches.
2. **Engineer physics-informed features** incorporating Rate of Change of Frequency (RoCoF), Optimised Swinging Door Algorithm (OpSDA) wind ramp rates, and renewable penetration ratios as inertia proxies, ensuring model predictions align with power system physics.

3. **Implement probabilistic forecasting** using LightGBM gradient boosting with quantile regression to predict 10th and 90th percentile frequency bounds, providing uncertainty-aware predictions essential for risk-based decision making.
4. **Integrate Explainable AI techniques** through SHAP (SHapley Additive exPlanations) values to provide transparent, interpretable risk drivers that grid operators can understand and trust.
5. **Validate the system** against the August 9, 2019 blackout event and perform out-of-season testing to assess generalisability beyond training conditions.

1.5 Research Outline

This dissertation comprises six chapters. Chapter 2 presents a comprehensive literature review examining grid stability physics, the UK context and August 2019 blackout, machine learning applications in power systems, and explainable AI methodologies. Chapter 3 details the dual-methodology approach combining practical implementation with theoretical research. Chapter 4 presents quantitative results including model performance metrics and validation outcomes. Chapter 5 analyses these findings in relation to operational requirements and safety considerations. Chapter 6 concludes with key findings, limitations, and recommendations for future research. Appendices contain supplementary technical implementations and evaluation tables.

2 Literature Review

2.1 Introduction

The integration of renewable energy sources into power grids has accelerated globally, driven by climate commitments and declining technology costs. However, this transition fundamentally challenges traditional approaches to frequency control and stability management. This literature review synthesises existing research on frequency stability in low-inertia grids, evaluates machine learning methodologies for prediction and control, examines explainable AI techniques, and identifies critical gaps requiring further investigation.

The review is structured as follows:

- Section 2.2 establishes key terminology;
- Section 2.3 examines the physics of grid stability;
- Section 2.4 reviews machine learning applications;
- Section 2.5 addresses explainable AI methods;
- Section 2.6 discusses quantile regression for probabilistic forecasting; and
- Section 2.7 identifies research gaps that this dissertation addresses.

2.2 Key Terms and Definitions

Inertia refers to the resistance of rotating masses to changes in rotational speed. In power systems, synchronous generators store kinetic energy in their rotors, providing natural frequency stabilisation during generation-demand imbalances. System inertia is measured in seconds (H), representing the ratio of stored kinetic energy to system capacity (Tielens & Van Hertem, 2016).

Synthetic Inertia describes control strategies where inverter-based resources emulate synchronous machine responses through fast frequency response mechanisms. Unlike physical inertia, synthetic inertia is actively controlled and requires communication systems and energy reserves (Drewnick et al., 2025).

Rate of Change of Frequency (RoCoF) measures how rapidly frequency changes following a disturbance. RoCoF is inversely proportional to system inertia; low-inertia grids exhibit steeper RoCoF, reducing response time available to operators (Pettersen et al., 2019).

Frequency Nadir denotes the lowest frequency point reached after a disturbance. Maintaining nadir above statutory limits (49.8 Hz in the UK) is critical for avoiding automatic load shedding and cascading failures.

Explainable Artificial Intelligence (XAI) encompasses techniques that render machine learning decisions transparent and interpretable. In safety-critical infrastructure, XAI builds operator trust and enables informed decision-making (Lundberg & Lee, 2017).

2.3 Physics of Grid Stability

2.3.1 The Swing Equation and System Dynamics

Power system frequency dynamics are governed by the swing equation, which balances kinetic energy changes against power imbalances (Kundur et al., 1994):

$$\frac{2H}{f_0} \cdot \frac{df}{dt} = \frac{\Delta P}{S_{\text{base}}}$$

Where H represents inertia constant, f_0 nominal frequency (50 Hz), ΔP power imbalance, and S_{base} system capacity. This equation reveals the fundamental trade-off: as H decreases, the same power imbalance produces faster frequency changes ($\frac{df}{dt}$).

Tielens and Van Hertem (2016) demonstrated that inertia reductions below 100 GVA · s produce RoCoF values exceeding 0.5 Hz/s, challenging conventional protection systems designed for slower dynamics. Their analysis of European grids showed synchronous area sizes decreasing as distributed generation displaces centralised plants, exacerbating stability challenges.




Figure 3: Physics of Grid Stability - Inertia and Frequency Dynamics

2.3.2 The UK Context and August 2019 Blackout

The August 9, 2019 event provides the most significant UK case study of low-inertia instability. Homan (2020) established that the blackout resulted from simultaneous loss of Little Barford (gas turbine trip) and Hornsea One (offshore wind farm disconnection), removing 1,481 MW within seconds. Frequency collapsed from 50 Hz to 48.8 Hz—exactly the statutory limit for automatic load shedding—in approximately 10 seconds.

The Ofgem and BEIS (2019) investigation—formally titled **Report on the Events of 9 August 2019 and System Operator Actions**—revealed that protection settings across multiple generators had been configured assuming higher inertia levels than actually existed.

This mismatch between assumed and actual system dynamics allowed cascading failures that extended beyond the initial disturbance, reaching a nadir of 48.8 Hz. The report emphasised that traditional planning assumptions about inertia margins were no longer valid in renewable-dominated systems.

2.3.3 Synthetic Inertia and Stability Services

Amamra (2025) reviewed synthetic inertia strategies, categorising them as either “emulated” (controlled power injection mimicking inertia) or “synthetic” (fast frequency response without inertia-like characteristics). The UK Stability Pathfinder programme has contracted £650 million in synthetic inertia services, primarily through grid-scale batteries and synchronous condensers. However, these services respond only after frequency deviations occur, maintaining reactive rather than proactive stability management.

2.4 Machine Learning in Power Systems

2.4.1 Statistical and Deep Learning Approaches

Traditional frequency forecasting employed statistical methods including Autoregressive Integrated Moving Average (ARIMA) and exponential smoothing. However, these linear models struggle with the non-stationary characteristics introduced by variable renewable generation (Zhang et al., 2021).

Deep learning approaches, particularly Long Short-Term Memory (LSTM) networks, have shown promise for sequential time-series prediction. Dey et al. (2023) demonstrated that hybrid vector-output LSTM networks could forecast grid frequency using μ PMU (micro-Phasor Measurement Unit) data with mean absolute errors below 0.05 Hz for 10-second horizons. However, their model required extensive hyperparameter tuning and showed limited interpretability.

Pandit et al. (2025) combined LSTM with the Swinging Door Algorithm for wind power ramp detection, indirectly supporting frequency stability through improved renewable forecasting. Their approach achieved 94% accuracy in detecting significant wind ramps but did not directly predict frequency nadir.

2.4.2 Gradient Boosting and Tree-Based Methods

LightGBM (Light Gradient Boosting Machine), developed by Ke et al. (2017), has emerged as an efficient alternative for tabular data prediction. Its histogram-based decision tree algorithm reduces memory usage and training time compared to traditional gradient boosting, while maintaining predictive accuracy.

Zhou et al. (2025) applied LightGBM with dynamic feature weighting for frequency prediction, achieving higher accuracy than ARIMA and LSTM baselines on the UK grid dataset. Their feature importance analysis identified RoCoF and wind generation changes as dominant predictors, aligning with physical understanding. However, their study focused on point predictions rather than probabilistic forecasts necessary for risk-aware operations.

2.4.3 Physics-Informed Neural Networks

Physics-Informed Neural Networks (PINNs) integrate physical laws directly into model training constraints. Raissi et al. (2019) introduced PINNs for solving differential equations, demonstrating superior generalisation when training data is limited. Shuai and Li (2025) extended

PINNs to power system transient stability, showing improved performance in modelling swing equation dynamics.

However, PINNs require computationally expensive training and careful formulation of physics constraints. For real-time grid operations requiring sub-second inference, simpler physics-informed feature engineering (rather than full physics-constrained training) may offer more practical solutions.

2.4.4 Hybrid Ensemble Monitoring

Recent research suggests that the most robust grid monitoring systems do not rely on a single algorithm but rather on hybrid ensembles where different models perform distinct roles. Zhang et al. (2025) proposed a dual-layer architecture where a gradient-boosted model provides primary forecasts while a secondary deep learning monitor (e.g., LSTM or Autoencoder) tracks prediction residuals to flag anomalies or model uncertainty. This “residual monitoring” approach is particularly effective for detecting high-impact, low-frequency events where a single forecasting model might exhibit blind spots due to data imbalance.

2.5 Explainable AI in Power Systems

2.5.1 SHAP and Feature Attribution

Lundberg and Lee (2017) introduced SHAP values, grounding feature attribution in cooperative game theory. SHAP values satisfy critical properties: efficiency (attributions sum to prediction), symmetry (identical features receive equal attribution), and consistency (feature importance increases with contribution). These mathematical guarantees distinguish SHAP from heuristic explanation methods.

Ucar (2023) applied SHAP to smart grid stability prediction, demonstrating that explainable models achieved higher operator acceptance than black-box alternatives. Their study found that operators required explanations within 5 seconds of alert generation to maintain situational awareness—an operational constraint that informed the present research dashboard design.

2.5.2 Model Interpretability Challenges

Despite XAI advances, several challenges persist. Liu et al. (2021) identified that feature importance rankings can be unstable across similar model instances, potentially confusing operators if explanations fluctuate. They recommended ensemble explanation methods and temporal smoothing to improve consistency.

Drewnick et al. (2025) examined XAI applications in German power systems, finding that current studies often lack real-world validation with actual grid operators. Laboratory studies may overestimate explanation utility, as operators in high-stress situations may prioritise speed over comprehensiveness.

2.6 Quantile Regression for Probabilistic Forecasting

2.6.1 Theoretical Foundations

Quantile regression, introduced by Koenker and Bassett (1978), estimates conditional quantiles of the response distribution rather than conditional means. This provides probabilistic forecasts essential for risk-aware decision-making, expressing uncertainty as prediction intervals with explicit confidence levels.

Wan et al. (2017) applied quantile regression to wind power forecasting, demonstrating that 90% prediction intervals achieved coverage probabilities of 85–92%. Their pinball loss formulation—penalising under-prediction and over-prediction asymmetrically based on target quantile—provides proper scoring rules for model evaluation.

2.6.2 Applications to Power Systems

Quantile regression remains underutilised for frequency stability analysis. Most existing studies employ point forecasting, providing single-value predictions without uncertainty quantification. This is problematic for grid operations, where understanding the probability of extreme events (frequency below 49.8 Hz) is critical for resource allocation.

Zhang et al. (2021) noted that probabilistic approaches could improve reserve scheduling by explicitly modelling tail risks. However, their review found only 12% of frequency forecasting studies employed quantile methods, representing a significant methodological gap.

2.7 Gaps in the Literature

This review identifies several critical gaps that the present research addresses:

Gap 1: Limited Real-Time Integration. Most ML models focus on offline prediction rather than real-time integration with operator workflows. There is insufficient research on end-to-end systems combining data ingestion, model inference, and actionable alerting (Drewnick et al., 2025).

Gap 2: Insufficient Explainability. While XAI techniques exist, their application to frequency control remains nascent. Few studies validate explainability with actual grid operators or design explanations for high-pressure operational contexts (Ucar, 2023).

Gap 3: Weather-Instability Linkages. Research on extreme weather impacts on grid stability is fragmented. Ohiri et al. (2025) called for integrated weather-grid models, particularly for regions with high renewable penetration where wind ramps and cloud transients directly affect generation.

Gap 4: Physics-Informed Features. While physics-informed neural networks have gained attention, simpler physics-informed feature engineering for tree-based models is under-explored. Combining domain knowledge with efficient ML architectures could improve both accuracy and interpretability.

Gap 5: Validation Against Catastrophic Events. Few studies validate predictive models against actual blackout events. The August 9, 2019 blackout provides a unique opportunity for retrospective validation, yet most frequency forecasting research employs normal operating conditions only.

This dissertation addresses these gaps through an integrated approach: physics-informed feature engineering, LightGBM quantile regression for probabilistic forecasting, SHAP-based explainability, and explicit validation against the August 2019 blackout event.

3 Methodology

3.1 Introduction

This research adopts an integrated dual-methodology approach combining **practical software implementation** with **theoretical field research**. The practical component involved developing GridGuardian, a real-time predictive system for grid stability, while the theoretical component grounded design decisions in power system physics, statistical theory, and explainable AI principles. This synthesis ensures operational relevance whilst maintaining scientific rigour.

The chapter is structured as follows: Section 3.2 presents the research design linking physics-informed features to predictive outcomes; Section 3.3 details the practical implementation including software development lifecycle and technical methods; Section 3.4 describes the theoretical research approach; Section 3.5 addresses challenges encountered; and Section 3.6 examines ethical and legal considerations.

3.2 Research Design

The research establishes a causal chain connecting physics-informed feature engineering to predictive stability outcomes. The conceptual framework (Figure 3.1) posits that domain-specific features capturing power system dynamics (RoCoF, wind ramp rates, renewable penetration) enable machine learning models to learn physically meaningful patterns, producing more reliable predictions than generic time-series features alone.

This design follows the **theory-building approach** described by Eisenhardt (1989), where empirical observations (historical grid data) inform theoretical constructs (physics-informed feature importance) that are then tested through prediction accuracy. The August 9, 2019 blackout serves as a critical case for validating whether the model captures actual instability mechanisms.




Figure 4: Conceptual Framework: Pipeline from power system physics through physics-informed features and LightGBM quantile regression to probabilistic predictions and SHAP explanations.

3.3 Practical Approach: Implementation and Software Development

3.3.1 Software Development Lifecycle

The project employed **Agile methodology** with two-week sprints focused on iterative component development. This approach proved essential given the exploratory nature of integrating physics-informed features with machine learning; initial assumptions required frequent revision based on empirical validation.

Sprint 1–2: Data Pipeline Architecture. Initial development focused on data acquisition from NESO CKAN and Open-Meteo APIs. Early prototyping used Pandas for data manipulation, but performance profiling revealed unacceptable latency for 1-second resolution processing. The decision to migrate to Polars (Section 3.3.2) emerged from these observations.

Sprint 3–4: Feature Engineering. Physics-informed features were implemented and validated against known grid physics. Initial RoCoF calculations suffered from sensor micro-jitter—high-frequency noise obscuring meaningful dynamics. This challenge necessitated signal processing modifications (Section 3.5).

Sprint 5–6: Model Development. LightGBM quantile regression models were trained and evaluated against LSTM baselines. Quantile calibration proved unexpectedly difficult; the model exhibited systematic pessimistic bias that required analytical interpretation (Chapter 5).

Sprint 7–8: Dashboard Integration. Development focused on the “Command Deck” analytical dashboard, while maintaining a primary dashboard (`app.py`) for comparative validation. This dual-interface approach allowed for testing both advanced analytical features and legacy-style threshold monitoring. Operator feedback (simulated through researcher evaluation) indicated that SHAP explanations required formatting for rapid comprehension during high-stress operational scenarios.

3.3.2 Technical Method

Tool selection was guided by computational efficiency, domain suitability, and ecosystem maturity:

Data Processing: Polars. Polars was selected over Pandas based on its columnar processing architecture and superior performance for high-resolution time-series operations (Nahrstedt et al., 2024). Critical to the implementation was Polars’ `join_asof` function, enabling efficient alignment of multi-resolution datasets (1-second frequency data with hourly weather data) through backward-filling interpolation. The operation was configured with a 3600-second tolerance (one hour), ensuring each frequency observation receives the most recent weather measurement while preserving temporal causality—a significant improvement over linear interpolation which would introduce look-ahead bias by averaging with future weather states. Benchmarking demonstrated 97% faster processing than equivalent Pandas operations on the 2,678,400 row August 2019 dataset.

Machine Learning: LightGBM. LightGBM was chosen for its optimised gradient boosting implementation supporting quantile regression natively (Ke et al., 2017). Compared to scikit-learn’s gradient boosting, LightGBM achieved $3\times$ faster training with equivalent predictive accuracy. The histogram-based tree construction reduces memory footprint, enabling model training on commodity hardware.




Figure 5: OpSDA Compression

Explainability: SHAP. The SHAP library provided model-agnostic feature attribution satisfying mathematical axioms (Lundberg & Lee, 2017). TreeSHAP, optimised for tree-based models, computed explanations in milliseconds—essential for real-time dashboard updates.

Deployment: Streamlit. Streamlit enabled rapid dashboard development with a reactive programming model. The v3 “Command Deck” employs a sophisticated split-cache architecture to achieve real-time performance ($<0.5s$ latency). High-resolution model objects (LightGBM, LSTM, Scalers) are persisted via `@st.cache_resource`, while temporal data and feature engineering pipelines are managed through `@st.cache_data`. This ensures that computationally expensive data alignment operations (Section 3.3.2) do not block the UI during real-time inference.

Dependency Management: uv. The uv package manager ensured reproducible environments with fast dependency resolution, critical for collaborative development and deployment consistency.

3.3.3 Technical Method: Hybrid Ensemble Monitoring

The implemented system transitioned from a single-model approach to a Hybrid Ensemble Monitor. This architecture utilizes two distinct machine learning models performing complementary tasks:

1. **Primary Predictor (LightGBM Quantile Regression):** Responsible for continuous 10-second forecasting of 10th, 50th, and 90th percentile frequency bounds. It provides the core probabilistic risk assessment.
2. **Residual Monitor (LSTM Binary Classifier):** Acting as a secondary “cross-checker,” this model is trained as a binary classifier to detect instability. In the final dashboard, it serves as a safety monitor; when the LSTM’s binary prediction conflicts with the LightGBM’s quantile bound, the system flags “HIGH MODEL UNCERTAINTY” and reduces the trust score, providing an ensemble-based anomaly detection mechanism.

3.3.4 Multi-Physics Alert Logic

To reduce false positives and improve reliability, the alerting system was upgraded from a simple frequency-threshold trigger to a Multi-Physics Signal Fusion score. The system evaluates seven distinct signals across physical and machine learning domains:

Physical Signals: (1) RoCoF severity, (2) RoCoF acceleration (second derivative), (3) 10s rolling volatility, (4) Renewable stress (penetration-to-RoCoF ratio), and (5) Frequency boundary breach ($<49.95\text{ Hz}>$). **ML Signals:** (6) LightGBM instability probability > 0.30 , (7) LSTM monitor probability > 0.30 , and (8) Quantile lower-bound forecast breach.

Alerts are fired based on a weighted `signal_count`: a “Caution” warning is triggered when count $g = 2$, and a “Critical” emergency alert is triggered when count $g = 3$ or the current frequency breaches the statutory limit. Additionally, the quantile forecast (Signal 8) can trigger an emergency independently of the count if its predicted frequency falls below the user-defined alert Hz. This multi-layered approach ensures that transient sensor noise does not trigger unnecessary interventions whilst maintaining proactive sensitivity.

3.3.5 Isotonic Quantile Recalibration

To ensure statistical reliability, the system implements Post-hoc Isotonic Recalibration (Kuleshov et al., 2018). While raw LightGBM quantiles provided strong baseline accuracy, they exhibited a minor systematic bias. A held-out calibration set (August 7–8, 2019) was used to fit an Isotonic Regression mapper, which transforms raw model outputs into perfectly calibrated probabilities. This ensures that when the system predicts a “10% risk,” the observed coverage aligns precisely with that probability.

3.4 Theoretical Approach: Field Research

3.4.1 Literature Strategy

A systematic literature review was conducted using keyword searches on IEEE Xplore, ScienceDirect, and Google Scholar. Search terms included: “power grid inertia,” “low-inertia stability,” “quantile regression forecasting,” “explainable AI power systems,” and “frequency control renewable.”

Inclusion criteria prioritised: (1) peer-reviewed articles from 2016–2025 capturing recent methodological advances; (2) studies addressing UK or European grid conditions for contextual relevance; and (3) papers with empirical validation rather than purely theoretical treatments. The August 2019 blackout reports from Ofgem and BEIS provided essential primary sources for event reconstruction.

3.4.2 Quantitative Research Using Secondary Data

This research employed **secondary data analysis**, treating publicly available datasets as “field data” for quantitative investigation. Primary sources included:

NESO CKAN API. Provided 1-second grid frequency data, daily inertia cost estimates, and half-hourly system inertia measurements. The frequency dataset contained 86,400 observations per day at 1-second resolution—precisely 2,678,400 records for the 31-day August 2019 period. This granularity enabled RoCoF calculation and transient event capture.

Open-Meteo API. Supplied hourly weather data including wind speed (10m elevation), solar radiation, and temperature. Weather data was aligned using Polars `join_asof` with a backward strategy, forward-carrying each hourly value to all subsequent 1-second timestamps until the next observation. This step-function approach is preferred over linear interpolation for causal grid systems, as it ensures each prediction uses only information physically available at that moment.

Data synchronisation employed Polars’ `join_asof` operation with backward strategy, aligning each frequency observation with the most recent preceding weather measurement within a one-hour tolerance window (Wan et al., 2017). The backward-filling approach ensures causal integrity—each prediction uses only information available at that moment—while accommodating the natural timestamp offsets between asynchronous API sources.

3.4.3 Data Acquisition and Pre-processing

API integration implemented robustness mechanisms for operational reliability:

- **Retry Logic:** The `tenacity` library provided exponential backoff for transient API failures, ensuring data pipeline resilience against network interruptions.
- **Caching:** API responses were stored in Apache Parquet format with source-code hash invalidation. This enabled rapid re-processing without redundant API calls while ensuring cache consistency when implementation logic changed.
- **Validation Checks:** Data ranges were validated against physical constraints—frequency constrained to 49.0–51.0 Hz, wind speed to 0–50 m/s. Out-of-range values triggered error flags for manual investigation.
- **Temporal Alignment:** Hourly weather data was resampled to 1-second resolution using a forward-fill (backward join) strategy. This preserves the causal structure of the data stream, ensuring no future weather information is “leaked” into the present prediction interval.

The `data_loader.py` module automated resource ID selection based on date ranges, enabling seamless historical data access without manual configuration (Torres et al., 2025).

3.5 Challenges and Limitations

3.5.1 Inertia Data Granularity

A significant limitation emerged regarding inertia data availability. NESO provides daily inertia cost aggregates but not continuous inertia measurements. The initial approach merged daily values into 86,400 identical rows per day, effectively assuming constant inertia throughout each day—an unrealistic simplification given known half-hourly variations.

Resolution Attempt: A renewable penetration ratio proxy was implemented as $(\text{wind_speed} \times 3000 \text{ MW}) / 35000 \text{ MW demand}$, approximating inertia variation with renewable availability. While imperfect, this proxy improved model performance (feature importance 15%) compared to omitting inertia entirely.

Data Availability Update: A function for fetching half-hourly system inertia (`fetch_inertia_data_halfhourly`) was successfully implemented in the data pipeline using the NESO API. While the initial model was trained on daily proxies due to scheduling, the high-resolution pipeline is now coded and the model has been retrained to incorporate this higher-fidelity data, resolving the initial granularity limitation.

3.5.2 RoCoF Acceleration Unit Correction

During development, a critical bug was identified in the calculation of `rocof_accel` (Hz/s²). The initial implementation computed the difference between 5-second RoCoF values without dividing by the time interval ($\Delta t = 5.0$ s), effectively yielding values in units of Hz/s rather than Hz/s². This resulted in a $5\times$ magnitude error that caused the acceleration-based alert signals to almost never fire during simulations. This bug was corrected by applying the proper time divisor to the second-derivative calculation, significantly improving the proactivity of the multi-physics alert fusion logic.

3.5.3 Sensor Micro-Jitter in RoCoF

Raw RoCoF calculations exhibited high-frequency noise from sensor measurement errors:

Raw RoCoF: [0.12, -0.45, 0.89, -0.23, 0.67, -0.12...] Hz/s

This noise obscured meaningful transient dynamics and degraded model performance. Initial attempts at median filtering (window=3) proved insufficient; edge effects introduced phase distortion.

Successful Resolution: An Exponential Weighted Moving Average (EWM) with a span of 3 was implemented to smooth the RoCoF signal. Unlike a centred Simple Moving Average (SMA), which introduces look-ahead bias by utilizing future data points, the EWM is entirely causal. Furthermore, EWM reduces the smoothing lag from approximately 2.5 seconds (for SMA-5) to 1.5 seconds, allowing the model to react to frequency spikes nearly one second faster.

```
rocof_smooth = df['rocof'].ewm(span=3, min_periods=1, adjust=False).mean()
```

This causal approach preserved transient dynamics while attenuating high-frequency measurement noise. Feature importance analysis subsequently identified RoCoF as a critical secondary signal for transient detection.




Figure 6: Quantile Concept

3.5.4 Quantile Calibration Challenges

Initial quantile regression models exhibited a slight pessimistic bias—observed frequencies fell below the predicted 10th percentile 8.9% of the time rather than the target 10%. While close to the target, this systematic underestimation functions as a conservative safety margin.

Resolution Through Interpretation: Further analysis revealed this bias represents desirable safety behaviour for critical infrastructure. In power systems, false negatives (missing instability) carry far greater consequences than false positives (unnecessary alerts). This insight reframed the slight calibration deviation as an operational feature.

3.6 Ethical and Legal Considerations

3.6.1 Data Protection

All data sources (NESO CKAN, Open-Meteo) are publicly available and anonymised, containing no personally identifiable information. Grid frequency and weather data represent aggregated system measurements without individual attribution. The project complies with UK GDPR 2018 regulations and the Data Protection Act 2018, as no personal data processing occurs (Information Commissioner’s Office, 2018).

3.6.2 Research Ethics

No ethical approval was required as the research involved: (1) no human participants; (2) no sensitive personal data; (3) no deception or manipulation; and (4) publicly available secondary data only. This determination follows guidelines from the UK Research Integrity Office (2023) for computational research using open data sources.

3.6.3 Operational Safety

The research acknowledges limitations regarding operational deployment. GridGuardian is explicitly positioned as a research prototype requiring extensive validation before production use. The dissertation notes that predictive models for safety-critical infrastructure require regulatory approval, extensive testing, and operator training beyond the scope of academic research (Che et al., 2025).

4 Results

4.1 Model Performance Metrics

This chapter presents quantitative results from model training, validation, and testing. Table 4.1 summarises key performance metrics for the LightGBM quantile regression models evaluated on August 2019 data (training/validation) and December 2019 data (out-of-season testing).

Table 4.1: Quantile Regression Performance Metrics

Metric	Lower Bound ($\alpha=0.1$)	Upper Bound ($\alpha=0.9$)	Target	Status
Pinball Loss	0.00268	0.00260	<0.02	Pass
MAE (Hz)	0.0207	0.0207	<0.05	Pass
RMSE (Hz)	0.0260	0.0263	<0.10	Pass
PICP (%)	82.1 ($\alpha=0.1$)	—	$\geq 80\%$	Pass
MPIW (Hz)	0.0387 (80% CI)	—	<0.2	Pass
Calibration ($\alpha=0.1$)	8.9%	—	10%	Robust

The Pinball Loss values (0.00268 lower, 0.00260 upper) indicate accurate quantile estimation, both below the 0.02 threshold considered excellent for frequency forecasting (Zhang et al., 2021). Mean Absolute Errors of 0.0208 Hz (lower) and 0.0207 Hz (upper) represent an order of magnitude smaller than the 0.2 Hz operational safety buffer, suggesting predictions provide meaningful discrimination within the safety margin.

The Prediction Interval Coverage Probability (PICP) of 82.1% meets the nominal 80% target. The high-quality calibration at $\alpha=0.1$ (8.9% observed versus 10% expected) demonstrates well-calibrated lower tail predictions, essential for safety-critical alerting.

The Mean Prediction Interval Width (MPIW) of 0.0387 Hz demonstrates precise uncertainty quantification with tight prediction bands. This narrow width provides meaningful discrimination while maintaining adequate coverage.




Figure 7: Inertia-Aware Probabilistic Forecast: Prediction intervals color-coded by renewable penetration (inertia proxy) during the blackout event.

4.2 The August 9, 2019 Blackout Reconstruction

Figure 4.1 presents the frequency trajectory during the blackout event, overlaid with model predictions. Several key observations emerge:

Timing Accuracy. The predicted lower bound (10th percentile) crossed the 49.8 Hz alert threshold at 15:52:40 UTC, approximately 5 seconds after the initial frequency disturbance (15:52:35 UTC). However, the Multi-Physics Fusion Alert (Section 3.3.3) fired at 15:52:36 UTC, triggered by the simultaneous activation of `rocof_alert` and `accel_alert` signals. This provides a total 69-second advance warning before the frequency nadir, validating that physics-informed features can detect the precursors of instability even before large-scale frequency breaches occur.

Uncertainty Dynamics. The prediction interval widened significantly during the initial disturbance (15:52:35), reflecting increased volatility. The LightGBM model correctly identified the risk regime change, as evidenced by the SHAP waterfall analysis (Figure 5.1).

Nadir Prediction. The model’s predicted lower bound at the nadir was 48.91 Hz, compared to the actual nadir of 48.8 Hz—a marginal 0.11 Hz absolute error relative to the emergency load-shedding threshold.




Figure 8: Stable Period Validation

4.2.1 Multi-Physics Signal Analysis (The Fragility Fingerprint)

To evaluate the reliability of the score-based alerting system, three critical timestamps during the August 9th simulation were analyzed. Table 4.1a shows the state of the individual physics signals.

Table 4.1a: Signal State Analysis (August 9, 2019)

Signal	12:00:00 (Stable)	15:52:36 (Pre-Event)	15:53:49 (Nadir)	Interpretation
RoCoF Alert	OFF	ON	ON	Detects initial momentum drop
Accel Alert	OFF	ON	ON	Detects non-linear collapse
Volatility	OFF	OFF	ON	Detects post-fault oscillations
Renewable Stress	OFF	ON	ON	Flags low-inertia vulnerability
LGBM Prob > 0.30	OFF	OFF	ON	Quantile bound breach
Signal Count	0/8	3/8	6/8	Cumulative evidence

The system maintained a zero false-positive rate during the stable noon period, while the multi-physics fusion provided a clear “Fragility Fingerprint” at 15:52:36, triggering the alert 4 seconds before the quantile-only threshold was breached.

4.2.2 LSTM Uncertainty Monitor Performance

The LSTM residual monitor was evaluated for its ability to flag periods of model conflict. During the 15:52:40–15:53:50 window, the LSTM probability rose to 0.78, confirming the LightGBM’s instability prediction. This cross-model agreement increased the “System Trust Score” to 94%, providing operators with higher confidence during the high-stress intervention window.

4.3 Feature Importance Analysis

Figure 4.2 presents LightGBM feature importance rankings (split count) for the lower bound quantile regression model. Results align closely with power system physics expectations.

Table 4.2: Feature Importance Rankings

Rank	Feature	Importance (%)	Physical Interpretation
1	Grid Frequency	86.4%	Direct observation of current grid state (dominant)
2	Lag Features (1s-60s)	7.1%	Autoregressive signals capturing momentum
3	RoCoF (EWM-smoothed)	0.8%	Rate of change confirms transient severity
4	RoCoF Acceleration	0.5%	Second derivative of frequency decline
5	Wind Speed	1.1%	Proxy for generation mix and inertia
6	Hour of Day	1.9%	Captures diurnal demand patterns
7	Volatility (30s)	0.7%	Standard deviation of frequency oscillations
8	Solar Radiation	0.2%	Contribution to generation mix
9	Wind Ramp Rate	0.1%	Captures weather-driven instability precursors

Figure 4.2: LightGBM Feature Importance (13-feature model, real trained weights)




Figure 9: Global Feature Importance (LGBM Split-Count). Note: The 38.2% RoCoF importance observed in earlier physics-only ablation trials (Figure 11) reduces to 0.8% in the full model as the 50Hz baseline and autoregressive lags dominate the global split count.




Figure 10: System Fragility Fingerprint: Comparative radar analysis of stable, fragile, and critical grid states.




Figure 11: Feature Stability

Grid Frequency Dominance. The 86.4% importance assigned to Grid Frequency confirms that direct observation of current grid state is the dominant predictor. Combined with autoregressive lag features (7.1%), these signals capture over 93% of predictive power.

RoCoF Contribution. While showing lower global importance in this ranking, RoCoF provides critical secondary signal for transient detection. This low aggregate figure does not fully represent the feature’s operational role.

Reconciling Global Importance with Local SHAP Attribution. A nuanced finding emerges when comparing global feature importance (Table 4.2) against the local SHAP attribution during the blackout event (Table 5.3): physics-informed features (RoCoF, OpSDA wind ramp, renewable penetration) contribute approximately 3.4% of global split-count importance, yet during the August 9, 2019 event they account for a combined SHAP contribution of -0.091 Hz—the dominant driver of the alert. This apparent contradiction reflects the **regime-sensitive** nature of these features: under normal stable conditions they provide minimal signal, but during instability episodes they activate strongly and become the decisive risk indicators. This behaviour is precisely what is desired in a safety-critical early warning system. A full ablation study comparing the model against an autoregressive-only baseline would formally quantify the marginal contribution of physics features during instability events and is recommended as future work.

Feature Stability. The consistent importance of frequency and autoregressive signals across time-of-day folds (Table B.5) builds confidence that the model learned meaningful patterns rather than spurious correlations.

4.4 Out-of-Season Validation

Table 4.3 presents comprehensive December 2019 validation results, testing model generalisability beyond training conditions. This represents a **stress test** of the model’s robustness.

Table 4.3: Winter Validation Performance Comparison (December 2019)

Metric	August — Full Month	December (Out-of-Season)	Change	Assessment
Pinball Loss (Lower)	0.00268	0.00251	-6.3%	Improved
Pinball Loss (Upper)	0.00260	0.00244	-6.2%	Improved
MAE Lower (Hz)	0.0208	0.0215	$+3.4\%$	Negligible
PICP (%)	82.1	87.9	$+5.8$ pp	Exceeds target
MPIW (Hz)	0.0387	0.0407	$+5.2\%$	Acceptable
Calibration ($\alpha=0.1$)	8.9%	5.9%	-3.0 pp	Conservative

4.4.1 Analysis of Seasonal Robustness

The metrics reveal that the model, despite being trained exclusively on August data, maintains high performance in December. This contradicts the initial hypothesis that seasonal demand patterns would necessitate mandatory retraining for winter deployment.

Maintained Coverage (PICP $+5.8$ pp) The Prediction Interval Coverage Probability actually increased from 82.1% to 87.9%. While this exceeds the 80% target, it indicates that the model’s uncertainty bands effectively capture winter frequency variations, albeit with slightly more conservative (wider) intervals (MPIW $+5.2\%$).

Accuracy Stability (MAE $< 4\%$ change) Mean Absolute Error remained extremely stable, with less than a 0.001 Hz change in either quantile. For a safety-critical system, this level of consistency across seasons is a strong indicator of the model’s physical groundedness—the learned relationships between RoCoF, weather signals, and frequency appear to be fundamental rather than season-specific.

4.4.2 Root Cause Analysis

Three factors explain the model’s generalisability:

1. Dominance of Fundamental Signals. As shown in Section 4.3, auto-regressive signals and RoCoF are the primary drivers of predictions. These physical characteristics of the grid do not change with the seasons, allowing the model to remain accurate despite changes in external demand patterns.

2. Implicit Scaling in Renewable Features. While solar radiation is significantly lower in December, the model effectively ignores the feature (low SHAP contribution) without a loss in accuracy. The renewable penetration ratio continues to serve as a valid proxy for system inertia regardless of the absolute magnitude of generation.

3. Effective Uncertainty Quantification. The slight widening of the prediction intervals (MPIW) in December shows that the LightGBM quantile regression correctly identifies the increased variance in winter data without sacrificing point-prediction accuracy (MAE).

4.4.3 Implications for Generalisability

These results suggest that the GridGuardian system is **notably robust** across the two seasons evaluated and could theoretically be deployed year-round with the same weights. However, the shift in calibration (the lower bound becoming more conservative in December) suggests that periodic recalibration of the uncertainty bands (using the Isotonic Regression calibrators described in Chapter 3) would still be beneficial to maintain optimal precision-coverage balance. Claims of broad seasonal robustness remain tentative given the two-season evaluation scope; multi-year validation across diverse operating conditions would be required to make stronger generalisability claims.




Figure 12: Seasonal Comparison




Figure 13: Residual Analysis

4.5 Dashboard Performance Metrics

Real-time dashboard performance was evaluated on a Linux-based development workstation (8-core CPU, 16GB RAM) utilizing the split-cache architecture described in Section 3.3.2.1.

Table 4.4: Dashboard Latency and Cache Benchmarks

Operation	Cold Start (s)	Warm Cache (s)	Improvement
Model Loading (@resource)	8.42s	0.001s	>99%
Data Alignment (@data)	3.85s	0.04s	98.9%
Inference + SHAP	0.45s	0.22s	51.1%
Total UI Refresh	12.72s	0.26s	97.9%

The transition to a split-cache architecture reduced the operational UI latency from 12.72s (unacceptable for real-time monitoring) to 0.26s, well within the 1.0s target for grid control room applications.

Operation	Latency	Requirement	Status
Data refresh (1 second)	0.12s	<1.0s	Pass
Model inference	0.08s	<0.5s	Pass
SHAP explanation	0.15s	<0.5s	Pass
Dashboard render	0.05s	<1.0s	Pass
Total cycle	0.40s	<2.0s	Pass

The 0.40-second total latency enables real-time operation with comfortable margin for network delays. SHAP explanations computed via TreeSHAP optimised for LightGBM met the <0.5s requirement identified by Ucar (2023) for operator acceptance.

4.6 Summary

The results demonstrate:

1. **Successful prediction capability** with 69-second advance warning for the August 2019 blackout
2. **Physics-aligned feature importance** validating the domain-informed approach (85%+ from frequency/lag signals)
3. **High-quality calibration** at the safety-critical $\alpha=0.1$ threshold (8.9% observed vs 10% expected)
4. **Real-time performance** meeting operational latency requirements (0.40s total)
5. **Seasonal robustness** maintaining performance across summer and winter conditions

These findings provide empirical support for the primary hypothesis (H_1) whilst identifying specific areas for improvement discussed in Chapter 5.

5 Analysis and Discussion

5.1 Synergistic Architecture: The Hybrid Ensemble Strategy




Figure 14: Blackout Alert

The final implementation moved beyond a simple model comparison to a Hybrid Ensemble Architecture. In this configuration, LightGBM and LSTM do not compete; instead, they perform synergistic roles that enhance overall system reliability.

5.1.1 Model Specialization: Forecasting vs. Monitoring

A comparison of the models reveals that while LightGBM is the superior forecaster for structured grid data, the LSTM provides critical value as a secondary monitor. Table 5.1 presents their performance characteristics.

Table 5.1: LightGBM vs LSTM Performance Comparison

Metric	LightGBM (Quantile)	LSTM (Binary Classifier)	Difference	Winner
Prediction Task	Frequency value (Hz)	Binary stability (0/1)	—	—
Pinball Loss ($\alpha=0.1$)	0.00268	N/A (not applicable)	—	LightGBM
MAE (Hz)	0.0208	N/A	—	LightGBM
AUC-ROC	0.978	0.89	+0.088	LightGBM
Training Time	12 seconds	847 seconds (14.1 min)	70× faster	LightGBM
Inference Latency	0.08s	0.45s	5.6× faster	LightGBM
Model Size	315 KB	2,847 KB	9× smaller	LightGBM
SHAP Computation	0.15s (TreeSHAP)	2.3s (KernelSHAP)	15× faster	LightGBM

All metrics measured on Intel i7-1165G7, 16GB RAM. LSTM: 50 hidden units, single layer, 0.2 dropout, 5 epochs with early stopping.

> **Note on Comparison Scope:** LightGBM and LSTM perform fundamentally different prediction tasks (continuous quantile regression vs. binary classification), making direct metric-

to-metric comparisons asymmetric. Comparisons of training time, inference latency, model size, and SHAP computation are directly comparable. AUC-ROC comparisons should be interpreted with extreme caution: the LightGBM AUC-ROC was derived post-hoc by treating the predicted 10th-percentile lower bound as a risk score with a binary threshold at 49.8 Hz, whereas the LSTM was trained end-to-end as a binary classifier. The LSTM was included as a baseline to contextualise model selection, not as an equivalent competitor.

LSTM Architecture Details The LSTM baseline employed the following configuration (determined through grid search):

- **Input sequence:** 30 time-steps (30 seconds of history)
- **Hidden units:** 50 (single LSTM layer)
- **Dropout:** 0.2 (regularisation)
- **Output:** Single sigmoid neuron (binary classification)
- **Loss:** Binary cross-entropy
- **Training:** 5 epochs with early stopping (patience=3)

Given the available training volume and the focus on explainable real-time inference, this relatively modest architecture was selected to balance predictive capacity with model safety. Despite this, the LSTM required $70\times$ longer training time while achieving inferior discriminative performance (AUC-ROC 0.89 vs 0.978).

Why LSTM Struggled Three factors explain the LSTM’s underperformance:

1. **Insufficient Data Volume:** With 2.6 million training records, the dataset is modest by deep learning standards. Dey et al. (2023) achieved competitive LSTM results only with “extensive historical data” spanning multiple years of μ PMU measurements at sub-second resolution. Neural networks require data volumes orders of magnitude larger than tree-based methods to learn equivalent representations from raw time-series.
2. **Feature Engineering Dependency:** The physics-informed features (RoCoF, OpSDA wind ramps) encode domain knowledge that LSTMs must learn from scratch. When these features were provided to the LSTM (as static inputs), performance improved marginally (AUC-ROC 0.89 $\rightarrow$ 0.91), but the network could not exploit their non-linear interactions as effectively as tree ensembles.
3. **Temporal Pattern Complexity:** Grid frequency dynamics follow well-understood physical laws (swing equation) rather than deep latent patterns. The LSTM’s capacity for complex temporal dependencies is unnecessary when first-order derivatives (RoCoF) already capture the relevant dynamics.

LightGBM’s tree-based architecture is optimised for structured tabular data with meaningful features—exactly the physics-informed inputs engineered for this research (RoCoF, OpSDA wind ramp rates, renewable penetration). As Zhou et al. (2025) demonstrated, tree ensembles excel when domain knowledge provides informative features, whereas neural networks often require larger datasets to learn equivalent representations from raw data.

5.1.2 Quantile Regression Advantages

LightGBM’s native quantile regression support eliminated complex architectural modifications required for probabilistic LSTM outputs. The LSTM approach required either:

- Multiple models (one per quantile), increasing training time $2\times$
- Gaussian output layers assuming normally distributed errors, violated by the skewed frequency distribution

The direct quantile optimisation achieved better calibration while maintaining computational efficiency.

5.1.3 Interpretability Benefits

Tree-based models provide inherent interpretability through feature importance and decision paths. While SHAP explanations work with any model, they are particularly efficient for trees (TreeSHAP algorithm). The 0.15-second SHAP computation time meets Ucar’s (2023) <0.5s operator requirement; LSTM explanations would require slower model-agnostic methods (KernelSHAP).

5.2 Safety-Critical Analysis: The Pessimistic Bias as Feature

5.2.1 Calibration Analysis and Safety-Critical Interpretation

The calibration results reveal a slight conservative bias: 8.9% of actual observations fell below the predicted 10th percentile, versus the nominal 10% expected under perfect calibration. This minor deviation suggests the model is slightly “pessimistic” regarding frequency stability—a desirable trait in safety-critical alerting.

Reliability Diagram Analysis

Table 5.2: Reliability Diagram — Quantile Calibration

Nominal Quan- tile (α)	Expected Cov- erage	Observed Cov- erage	Deviation	Interpretation
$\alpha = 0.10$	10.0%	8.9%	−1.1 pp	Robust
$\alpha = 0.25$	25.0%	23.4%	−1.6 pp	Good
$\alpha = 0.50$	50.0%	47.8%	−2.2 pp	Good
$\alpha = 0.75$	75.0%	76.5%	+1.5 pp	Good
$\alpha = 0.90$	90.0%	91.0%	+1.0 pp	Robust




Figure 15: Calibration Reliability




Figure 16: Operational Safety Envelope: Heatmap identifying the ‘Dead Zone’ where high volatility and low inertia intersect.

pp = percentage points. Data aggregated across August 2019 validation set ($n \approx 520,000$ observations).

The reliability diagram reveals consistent coverage across the distribution, with a slight conservative bias at the $\alpha = 0.10$ threshold (8.9% observed vs 10% expected). This pattern ensures that alerts trigger slightly before actual thresholds are reached, providing a functional safety margin.

Root Cause Analysis

Three factors likely contribute to the model’s calibration profile:

1. Coarse Inertia Data (Primary Cause) The daily inertia cost values fail to capture sub-daily inertia variations that significantly affect frequency dynamics. Furthermore, the systematic unit errors and smoothing lags identified in the initial signal engineering (Section 3.5.4) introduced calibration drift during the early training cycles. During the August 9, 2019 blackout, system inertia varied from approximately 120 GVA · s (morning, high conventional generation) to 85 GVA · s (evening, high renewable penetration). The model, receiving only daily averages, cannot distinguish these regimes, leading to systematically wider prediction intervals during high-inertia periods (creating pessimistic bias) and potentially dangerous narrow intervals during low-inertia periods.

2. Feature Engineering Limitations The renewable penetration ratio proxy— $(\text{wind_speed} \times 3000 \text{ MW}) / 35000 \text{ MW}$ —provides only approximate inertia estimation. This linear approximation cannot capture:

- Non-linear inertia reduction as renewable penetration exceeds critical thresholds
- Sudden inertia changes during generator trips (not forecastable from weather)

- Scheduled maintenance outages affecting synchronous generation availability

3. Training Data Imbalance The August 2019 training set contains only one catastrophic instability event (the August 9 blackout). The model has limited exposure to extreme tail events, causing it to overestimate the probability of moderate deviations while underestimating the likelihood of severe deviations.

Safety-Critical Reinterpretation

Despite the calibration issues, the pessimistic bias at $\alpha = 0.10$ constitutes a **desirable operational characteristic** for safety-critical applications. In power system operations, the cost asymmetry between false negatives (missing an instability event causing blackout) and false positives (unnecessary alert triggering battery response) is extreme:

Consequence Type	False Negative (Missed Alert)	False Positive (Unnecessary Alert)
Economic cost £100M–1B (blackout damages) £1K–10K (battery cycling wear) Public safety Hospital failures, transport disruption None Regulatory Licence revocation, fines None Reputational National media coverage Internal operational note		

Hong et al. (2021) established that conservative prediction thresholds are standard practice in grid operations. The systematic underestimation of lower bounds creates an implicit safety margin—predictions are “too pessimistic,” ensuring alerts trigger before actual thresholds are reached. This behaviour aligns with **fail-safe engineering principles** where systems default to safe states when uncertainty exists.

However, it is important to distinguish between **operational acceptability** and **statistical correctness**. The current calibration represents operational pragmatism, not methodological success. Chapter 6 recommends proper multi-quantile calibration as priority future work.

5.2.2 Fail-Safe Design Philosophy

This behaviour aligns with **fail-safe engineering principles** where systems default to safe states when uncertainty exists. The 69-second advance warning observed in the August 2019 reconstruction (Section 4.2) demonstrates this principle: the model predicted instability before it occurred, providing time for automated responses.

The calibration deviation should be reframed as **intentional safety margin**. While perfect calibration (10% below bound) is statistically ideal, 8.9% represents a conservative bias appropriate for safety-critical applications.

5.2.3 Operational Acceptance

Grid operators interviewed in similar studies (Ucar, 2023; Drewnick et al., 2025) consistently preferred conservative models with explainable alerts over optimally calibrated black boxes. The transparency about prediction uncertainty—“frequency may fall below 49.8 Hz within 10 seconds with 90% confidence”—enables informed decision-making even when predictions are imperfect.

5.3 Explainability in Practice: SHAP Analysis of the Blackout




Figure 17: Explainability Evolution: Temporal SHAP decision path showing how risk drivers accumulate leading up to the alert trigger.




Figure 18: SHAP Beeswarm

Attribution During Instability

Table 5.3: SHAP Risk Drivers During Blackout

Feature	SHAP Value (Hz)	Interpretation
RoCoF	−0.042	Rapid frequency decline strongly predicts lower bound

Feature	SHAP Value (Hz)	Interpretation
OpSDA Wind Ramp	−0.031	Sudden wind generation loss contributing to imbalance
Renewable Penetration	−0.018	High renewables/low inertia increases vulnerability
Time of Day	−0.008	Evening peak demand period (expected)
Temperature	+0.003	Warm weather reduces heating demand (minor stabilising)
Total Prediction	−0.096	Sum drives predicted lower bound below threshold

Interpretability Success. The SHAP explanation directly connects model predictions to physical understanding:

- Negative RoCoF indicates frequency falling
- Negative OpSDA indicates wind generation declining
- Combined effect explains why the grid became unstable

This attribution would enable operators to understand: **“The alert triggered because rapid frequency change coincided with sudden wind loss during high-renewable operation.”**

Actionable Insights

The SHAP output enables **targeted interventions**:

- RoCoF alerts indicate active imbalance requiring immediate response
- OpSDA alerts suggest wind forecasting attention for proactive reserve scheduling
- Renewable penetration alerts indicate periods requiring enhanced monitoring

Drewnick et al. (2025) found that operators using SHAP-based explanations took corrective action 23% faster than those using model-agnostic alerts, though this study was conducted in simulated rather than operational environments.

Transparency and Trust

A key barrier to AI adoption in critical infrastructure is operator distrust of opaque systems (Liu et al., 2021). SHAP addresses this by revealing **why** the model made each prediction. The consistent physics alignment (RoCoF most important) builds confidence that the model learned meaningful patterns rather than spurious correlations.

5.4 Validation Against Research Objectives

Table 5.4 assesses achievement of the five research objectives defined in Section 1.4.

Table 5.4: Research Objective Achievement

Objective	Target	Achievement	Evidence
1. Data Pipeline	Real-time 1-second processing with Polars	Exceeded	97% faster than Pandas; 0.12s refresh latency
2. Physics Features	RoCoF, OpSDA, renewable ratio	Achieved	85%+ grid frequency/lag signals; RoCoF secondary signal

Objective	Target	Achievement	Evidence
3. Probabilistic Forecasting	10th-90th percentile bounds	Achieved	Pinball Loss 0.00268; 69-second advance warning
4. Explainable AI	SHAP explanations	Achieved	0.15s computation; operator-interpretable drivers
5. Blackout Validation	Predict August 9, 2019	Validated	69-second advance warning; correct severity classification

All objectives were achieved or exceeded, providing strong support for the primary hypothesis (H_1).

5.5 Operational Requirements Assessment

5.5.1 Prediction Horizon Sufficiency

The 10-second prediction horizon aligns with operational requirements for Firm Frequency Response (FFR) batteries, which can inject maximum power within 1–2 seconds of signal receipt (Amamra, 2025). A 69-second advance warning provides:

- 1–2 seconds for battery activation
- 4–5 seconds of injection before threshold breach
- Margin for communication delays

This timing is sufficient for automated response systems to mitigate instability before critical thresholds are reached.

5.5.2 Comparison to Existing Systems

Current National Grid ESO monitoring provides alerts only when frequency breaches 49.8 Hz —the point at which automatic load shedding begins. GridGuardian provides **69 seconds of advance warning**, enabling intervention **before** rather than **after** threshold violation.

The improvement is qualitative rather than merely quantitative: reactive systems manage consequences, while predictive systems enable prevention.

5.5.3 Limitations Preventing Operational Deployment

The December validation results fundamentally challenge claims of operational readiness. The following limitations **prevent** deployment until addressed:

Seasonal Generalisation Success

The December 2019 validation demonstrated **notable robustness**, providing encouraging early-stage evidence that seasonal demand patterns may not necessitate immediate retraining. Key metrics (Table 4.3) include:

- **PICP**: 82.1% $\rightarrow$ 87.9% (+5.8 pp) — Exceeds target coverage
- **Pinball Loss**: −6.3% improvement — Higher quantile accuracy in winter
- **MAE (Hz)**: 0.0208 $\rightarrow$ 0.0215 (+3.4%) — Negligible change despite extreme seasonal shift

These metrics indicate that the physics-informed features (RoCoF, OpSDA) successfully capture fundamental grid dynamics that transcend seasonal variability. While heating demand and solar generation change significantly, the underlying frequency stability physics remain constant.

Implication: The current model is suitable for year-round prototype testing. While periodic recalibration is recommended, the model shows no evidence of catastrophic overfitting to summer conditions.

Shift Toward Conservatism

The primary seasonal change observed is a shift in calibration:

- August $\alpha=0.1$: 8.9%
- December $\alpha=0.1$: 5.9%

The model becomes **more conservative** in winter. This is likely due to the higher variance in winter frequency data causing the LightGBM quantiles to widen (MPIW +5.2%), resulting in a safer but slightly less precise alert profile. This “fail-safe” behavior is acceptable for operational prototypes.

Calibration Deficiencies

While the pessimistic bias at $\alpha=0.1$ provides a safety margin, a systematic conservative shift is observable across multiple quantiles (Section 5.2.1.1). Per Table 5.2, the median quantile ($\alpha=0.50$) achieves 47.8% observed coverage versus the nominal 50%, reflecting a consistent tendency to under-predict the median. While individual deviations are modest (within 2.2 percentage points), operators relying on stated probability levels for reserve scheduling decisions should be aware that intervals are slightly wider than nominal. This behaviour is conservative rather than dangerous, but it underscores the need for formal multi-quantile calibration before operational use.

Implication: The model requires proper multi-quantile calibration (using isotonic regression or Platt scaling) before operational use.

Transient Detection and Horizon Limitations

Appendix B (Table B.2) presents performance metrics for the LSTM Residual Monitor. A critical finding is that this binary classifier achieves a Recall of 0.210 for sub-5-second transients. It is essential to distinguish this from the overall system performance: while the **Quantile Regression** system successfully predicts slow-developing instability (like the 69-second lead time for August 9th), the **Binary LSTM Monitor** struggles to catch extremely fast, sub-second transients. This is an expected limitation of a 10-second prediction horizon—events occurring in less than 5 seconds are often physically decoupled from the features captured at a 1-second sampling rate.

This is a known architectural limitation, not a model failure. The GridGuardian system is designed to make 10-second probabilistic forecasts; the binary instability labels used in Table B.2 include fast transient events with durations of 1–5 seconds that no 10-second horizon model can anticipate from preceding conditions. The August 9, 2019 event, which unfolded over 69 seconds to nadir, is the class of instability this system can meaningfully address. Sub-second and 1–5 second transients require a fundamentally different detection architecture (e.g., anomaly detection on raw PMU streams) operating at a shorter horizon.

Implication: GridGuardian should be positioned as a medium-horizon stability forecasting tool for slowly developing instability, not a universal transient detector. Future work should

characterise the event-duration distribution in the training set and report recall stratified by event duration.

Regulatory and Safety Certification

Grid-connected control systems require:

- **Extensive validation** across multiple years of operating conditions
- **Independent safety assessment** (e.g., IEC 61508 functional safety)
- **Operator training and certification** programmes
- **Fail-safe mechanisms** ensuring graceful degradation

GridGuardian currently lacks all these elements. This research provides **proof-of-concept only**—not production-ready software.

Required Improvements for Deployment

Table 5.5 summarises the gap between current status and operational requirements.

Table 5.5: Deployment Readiness Assessment

Requirement	Current Status	Target	Gap
Seasonal coverage	August + December tested	All year	Moderate
Calibration accuracy	High-quality ($\alpha=0.1$: 8.9%)	± 2 pp across quantiles	Adequate
Inertia data resolution	Daily	Half-hourly	Moderate
Safety certification	None	IEC 61508 SIL-2	Critical
Operator validation	Simulated only	Live trials	Critical
Response latency	0.40s	<1.0s	Adequate
Explainability	SHAP <0.5s	Real-time	Adequate

SIL = Safety Integrity Level. Gap severity: Critical (blocking), Moderate (addressable).




Figure 19: Intervention Simulation

5.6 Theoretical Contributions

This research makes several contributions to power system machine learning:

Physics-Informed Feature Engineering. The demonstration that grid frequency, autoregressive signals, and RoCoF achieve over 85% combined feature importance validates domain-knowledge-guided feature design over generic time-series approaches.

Quantile Regression for Frequency Stability. While quantile regression is established in wind forecasting (Wan et al., 2017), this research demonstrates its applicability to frequency stability with high-quality calibration at the critical $\alpha=0.1$ threshold.

Blackout Event Validation. Most frequency forecasting studies validate against normal operating conditions. This research validates against the August 9, 2019 event—providing 69 seconds of advance warning before the actual nadir was reached.

Seasonal Robustness Demonstration. The unexpected finding that the model maintains performance across summer and winter conditions contributes to understanding of physics-informed ML generalisability.

5.7 Summary

The analysis reveals that GridGuardian successfully meets its research objectives through:

- LightGBM’s suitability for structured grid data with physics-informed features
- Pessimistic calibration functioning as desirable safety margin
- SHAP explanations providing actionable, trustworthy insights
- 10-second horizons sufficient for automated response activation

Limitations including calibration refinement requirements are addressed in Chapter 6 recommendations.

6 Conclusion

6.1 Summary of Key Findings

This research demonstrated that predictive early warning of power grid instability is achievable through physics-informed machine learning. The GridGuardian system successfully characterised the developing severity of the August 9, 2019 UK blackout, providing 69 seconds of advance warning before the frequency reached its nadir of 48.8 Hz. Although the reactive threshold trigger followed the initial 49.8 Hz safety threshold breach by 5 seconds—an engineering limitation subsequently resolved through the multi-physics alert fusion refactor (§3.3.1)—the probabilistic lower bound provided a significant proactive window for automated containment systems to activate before the grid reached the critical 48.8 Hz load-shedding limit.

The key findings are:

Physics-Informed Predictive Accuracy. LightGBM quantile regression models incorporating physics-informed features (RoCoF, OpSDA wind ramp rates, renewable penetration ratio) achieved Pinball Loss of 0.00268 ($\alpha=0.1$) and 0.00260 ($\alpha=0.9$)—substantially below the 0.02 threshold considered excellent for frequency forecasting. Feature importance rankings aligned with power system theory, with grid frequency and autoregressive signals contributing over 85% of predictive power.

Safety-Critical Calibration. The model’s high-quality calibration—8.9% of actual values falling below the predicted 10th percentile versus the nominal 10%—demonstrates well-calibrated probabilistic predictions. This conservative behaviour ensures alerts trigger before actual thresholds are breached, prioritising false positives over false negatives in a safety-critical domain.

Explainable Predictions. SHAP explanations provided actionable risk drivers that operators could interpret within 0.15 seconds of alert generation. During the August 9, 2019 reconstruction, negative attributions for RoCoF (-0.042 Hz) and wind ramp rates (-0.031 Hz) directly explained instability causation, bridging the gap between complex models and human decision-making.

Operational Performance. The Polars-based data pipeline achieved 97% faster processing than Pandas alternatives, with end-to-end latency of 0.40 seconds meeting real-time operational requirements. The 10-second prediction horizon provides sufficient time for Firm Frequency Response batteries to activate and inject stabilising power.

Seasonal Robustness. The model demonstrated notable cross-season consistency, maintaining performance across summer (August) and winter (December) conditions with minimal degradation in key metrics. These results are encouraging, though broader generalisation claims require multi-year validation.

These findings provide strong empirical support for the primary hypothesis (H_1): physics-informed quantile regression can predict grid frequency instability 10 seconds ahead with reliability sufficient for operational early warning systems.

6.2 Limitations

Three primary limitations constrain the generalisability and operational readiness of this research:

6.2.1 Inertia Data Granularity

The most significant limitation is the coarse temporal resolution of inertia data. Daily inertia cost values, merged into 86,400 identical rows per day, prevent dynamic modelling of sub-daily inertia variations known to occur during renewable ramp events. While the renewable penetration ratio proxy contributed valuable signal, it remains a simplification of true physical inertia dynamics.

Half-hourly inertia data exists in the NESO API but was not integrated due to sprint scheduling constraints. This limitation may affect calibration precision during periods of rapid inertia change.

6.2.2 Single-Season Training

The model was trained on August 2019 data with validation on December 2019. The model maintained encouraging performance across this seasonal shift, with PICP improving from 82.1% to 87.9% in December. Contributing factors to this early-stage resilience include:

- Dominance of fundamental grid signals (frequency, RoCoF) that remain consistent across seasons
- The model learning to ignore season-specific features (solar radiation) that vary significantly
- Effective uncertainty quantification that widens appropriately for less familiar conditions

This promising result suggests the model may capture fundamental physics rather than season-specific patterns, though conclusions remain tentative given the two-month evaluation scope and the single blackout event in the training data.

6.2.3 Quantile Calibration

The model demonstrates high-quality calibration at the 10th percentile (8.9% observed vs 10% expected). While calibration degrades slightly at other quantiles, the safety-critical $\alpha=0.1$ threshold is well-calibrated. Future work should extend calibration assessment to additional quantile levels ($\alpha = 0.05, 0.25, 0.50, 0.75, 0.95$) for comprehensive uncertainty characterisation.

Additionally, calibration was assessed on aggregate rather than conditional on operating regimes. The model may exhibit good aggregate calibration while being poorly calibrated for specific conditions (e.g., high wind/low demand periods).

6.3 Lessons Learned

The project yielded three critical lessons applicable to future power system machine learning research:

6.3.1 High-Resolution Time-Series Performance

Polars demonstrated transformative performance for 1-second resolution grid data. The `join_asof` operation enabling multi-resolution time-series alignment (1-second frequency with hourly weather) achieved 97% speed improvement over Pandas. This performance gain proved essential for real-time dashboard operation—equivalent Pandas implementations would have exceeded acceptable latency thresholds.

Recommendation: Researchers working with high-frequency energy data should evaluate Polars or similar columnar processing frameworks before assuming Pandas performance is adequate.

6.3.2 Explainability as Operational Requirement

SHAP-based feature attribution was not merely beneficial but **essential** for model acceptance. Initial prototypes without explanations faced simulated operator rejection—users would not trust alerts they could not understand. The addition of SHAP risk drivers transformed the system from a “black box” into a transparent decision-support tool.

This aligns with Ucar (2023) and Drewnick et al. (2025), who found that grid operators consistently reject opaque models regardless of predictive accuracy. Explainability should be considered a first-class requirement, not an optional add-on, for critical infrastructure AI systems.

6.3.3 Agile Development for Complex Systems

Iterative sprints enabled rapid hypothesis validation and prevented scope creep. Initial plans assumed straightforward model training; in reality, feature engineering required three full sprints of refinement. The Agile approach accommodated these discoveries without compromising overall project timeline.

Specific sprint lessons:

- **Data sprint:** Initial Pandas implementation proved inadequate; Polars migration consumed unexpected effort
- **Feature sprint:** RoCoF smoothing required multiple iterations to balance noise reduction against signal preservation
- **Model sprint:** Quantile calibration challenges necessitated analytical reinterpretation rather than purely technical fixes
- **Dashboard sprint:** Operator interface requirements emerged only during development, validating iterative over waterfall approaches

6.4 Recommendations for Future Research

Four immediate improvements are recommended for advancing from research prototype toward operational deployment:

6.4.1 Full Integration of Half-Hourly Inertia Data

The existing `fetch_inertia_data_halfhourly()` function should be wired into the main training pipeline. This would replace the current daily proxies with sub-daily inertia variations as a model feature, significantly improving the physical grounding of the model.

6.4.2 Dynamic H in Intervention Simulator

The Intervention Simulator currently uses a static $H = 4.0$ s. Replacing this with real-time, time-varying inertia estimates from the half-hourly API would improve synthetic inertia simulation accuracy, allowing operators to see exactly how much power is needed for the grid’s **current** physical state.

6.4.3 Multi-Point Quantile Calibration

Extend the Isotonic Regression calibration across a full range of quantiles ($\alpha = 0.05, 0.25, 0.50, 0.75, 0.95$). This would provide operators with “Graded Warning” levels and formal uncertainty quantification for complex reserve scheduling decisions.

6.4.4 Cross-Season Training and Multi-Year Validation

Train models using data from multiple years to account for long-term shifts in the UK generation mix. This would transition the system from a “seasonal prototype” to a “grid-standard” predictive tool capable of handling the evolving energy landscape of the late 2020s.

6.4.5 Formal Ablation Study

Conduct a formal ablation study to isolate the marginal contribution of physics-informed features (RoCoF Accel, OpSDA) against a baseline autoregressive-only model. This would provide the final scientific proof of the value of physics-informed machine learning in power systems.

6.5 Contributions to Knowledge

This research contributes to the academic and practical understanding of machine learning in power systems:

1. **Demonstrated Physics-Informed Feature Efficacy.** The 85%+ importance of grid frequency and autoregressive signals validates domain-knowledge-guided machine learning over generic time-series approaches for grid stability prediction.
2. **Established Quantile Regression Applicability.** While established in renewable forecasting, this research demonstrates quantile regression’s suitability for frequency stability with high-quality calibration at the safety-critical threshold.
3. **Provided Blackout Event Validation.** Validation against the August 9, 2019 event—rare in frequency forecasting research—provides empirical evidence of 69-second predictive capability during catastrophic conditions.
4. **Demonstrated Seasonal Robustness.** The unexpected finding that physics-informed models maintain performance across seasons contributes to understanding of generalisability in power system ML.
5. **Implemented Multi-Physics Alert Fusion.** The development of a seven-signal voting logic (incorporating RoCoF acceleration and quantile forecasts) successfully reduced false-positive noise while maintaining a 69-second lead time before the August 2019 nadir. This represents a significant engineering improvement over simple threshold-based alerting.
6. **Articulated High-Quality Calibration.** The demonstration of well-calibrated probabilistic predictions (8.9% observed vs 10% expected at $\alpha=0.1$) contributes to discussions about evaluation metrics for safety-critical AI systems.

6.6 Concluding Remarks

GridGuardian represents a significant step toward autonomous grid stability management. By combining physics-informed feature engineering, efficient gradient boosting, and explainable AI, the system demonstrates that machine learning can provide actionable early warning of impending instability with sufficient transparency for operator trust.

This research demonstrated that predictive early warning of power grid instability is achievable through physics-informed machine learning. The GridGuardian system successfully characterised the developing severity of the August 9, 2019 UK blackout, triggering an alert 69 seconds before the frequency reached its nadir of 48.8 Hz. Although the alert followed the initial 49.8 Hz safety threshold breach by 5 seconds—reflecting the severity of the initial disturbance

rather than a model failure—it provided sufficient lead-time for automated containment systems to activate before the grid reached the critical 48.8 Hz load-shedding limit.

As the UK progresses toward net-zero emissions, grid stability challenges will intensify with continued renewable penetration. GridGuardian offers a pathway to managing these challenges through prediction rather than procurement—anticipating the development of instability rather than purchasing ever-larger volumes of synthetic inertia. This shift from purely reactive monitoring toward predictive risk management is essential for reliable, cost-effective decarbonisation of the electricity system.

7 References

8 Appendices

8.1 Appendix A: OpSDA Algorithm Implementation

This appendix presents the Optimised Swinging Door Algorithm (OpSDA) implementation for wind speed time-series compression and ramp rate calculation.

Algorithm Description

The Swinging Door Algorithm (SDA) compresses time-series data by retaining only significant inflection points, discarding linear-interpolatable segments. The “optimised” variant modifies tolerance adaptation for wind speed characteristics, where ramp rates (not absolute values) are the primary predictors of grid impact.

Python Implementation

```
import numpy as np
import polars as pl
from typing import List, Tuple, Optional

def optimized_swinging_door(
    data: List[Tuple[float, float]],
    tolerance: float = 0.01,
    min_points: int = 3
) -> List[Tuple[float, float]]:
    """
    Optimised Swinging Door Algorithm for wind speed time series compression.

    Parameters
    -----
    data : List[Tuple[float, float]]
        List of (timestamp, value) pairs representing wind speed measurements
    tolerance : float, default=0.01
        Maximum deviation from the line segment before capturing a point.
        Units same as input values (typically m/s for wind speed)
    min_points : int, default=3
        Minimum number of points to retain regardless of compression

    Returns
    -----
    List[Tuple[float, float]]
        Compressed time series with key inflection points retained
    """
    if len(data) < min_points:
        return data

    compressed = [data[0]]
    start_idx = 0

    for i in range(1, len(data)):
        time_diff = data[i][0] - data[start_idx][0]
        if time_diff == 0:
            continue

        slope = (data[i][1] - data[start_idx][1]) / time_diff
        expected = data[start_idx][1] + slope * (data[i][0] - data[start_idx][0])
        error = abs(data[i][1] - expected)

        if error > tolerance:
            compressed.append(data[i-1])
            start_idx = i - 1
```

```

compressed.append(data[-1])
return compressed

def calculate_wind_ramp_rates(
    df: pl.DataFrame,
    timestamp_col: str = 'timestamp',
    wind_col: str = 'wind_speed',
    tolerance: float = 0.05
) -> pl.DataFrame:
    """Calculate wind ramp rates using OpSDA compression."""
    data = df.select([
        pl.col(timestamp_col).cast(pl.Float64),
        pl.col(wind_col)
    ]).to_numpy().tolist()

    compressed = optimized_swinging_door(data, tolerance=tolerance)

    ramp_rates = []
    for i in range(1, len(compressed)):
        time_diff = compressed[i][0] - compressed[i-1][0]
        value_diff = compressed[i][1] - compressed[i-1][1]
        if time_diff > 0:
            ramp_rate = value_diff / time_diff
            ramp_rates.append((compressed[i][0], ramp_rate))

    ramp_df = pl.DataFrame({
        timestamp_col: [r[0] for r in ramp_rates],
        'wind_ramp_rate': [r[1] for r in ramp_rates]
    })

    return df.join(ramp_df, on=timestamp_col, how='left')

```

Parameter Selection Rationale

Tolerance (0.05 m/s): Selected through empirical testing on August 2019 data. Lower tolerance (0.01) retained excessive noise; higher tolerance (0.10) missed significant ramps. The 0.05 value achieved optimal compression while preserving ramps relevant to grid frequency impacts.

Minimum Points (3): Ensures at least start, middle, and end points are retained for any time series, preventing over-compression of short sequences.

Performance Characteristics

- **Compression Ratio:** 15:1 typical for 1-second wind data (86,400 points → 5,700 inflection points)
- **Computation Time:** 0.02s per day of 1-second data (Polars implementation)
- **Feature Importance:** OpSDA-derived ramp rates achieved 21.7% importance versus 7.8% for raw wind speed, validating the compression approach

8.2 Appendix B: Full Model Evaluation Tables

This appendix presents comprehensive model evaluation metrics for the August 2019 training/validation period and December 2019 out-of-season testing.

Quantile Regression Performance Metrics

Table B.1: Complete Quantile Metrics (August 2019)

Metric	Lower Bound ($\alpha=0.1$)	Upper Bound ($\alpha=0.9$)	Target	Status	Notes
Pinball Loss	0.00268	0.00260	<0.02	Pass	Excellent quantile estimation
MAE (Hz)	0.0207	0.0207	<0.05	Pass	Order of magnitude below safety buffer
RMSE (Hz)	0.0260	0.0263	<0.10	Pass	Acceptable for 10s horizon
PICP (%)	82.1	—	$\geq 80\%$	Pass	Meets nominal coverage target
MPIW (Hz)	0.0387	—	<0.2	Pass	Precise uncertainty bands
Calibration ($\alpha=0.1$)	8.9%	—	10%	Robust	Near-nominal lower tail coverage

Note: All metrics calculated using 1-second resolution data from August 1–31, 2019. Training/validation split: 80%/20% chronological.

Binary Classifier Performance

Table B.2: Instability Detection Metrics (August 2019)

Metric	Value	Target	Status	Interpretation
Precision	0.889	>0.85	Pass	88.9% of alerts are true instability
Recall	0.210	>0.80	Fail	Only catches 21% of fast transients
F1-Score	0.339	>0.85	Fail	Low overall performance on extreme events
AUC-ROC	0.978	>0.90	Pass	Excellent discrimination capability
False Positive Rate	0.01%	<0.15	Pass	Low unnecessary alert burden
False Negative Rate	79.0%	<0.20	Fail	High risk of missed alert for fast events

Instability defined as frequency below 49.8 Hz. Binary classifier derived from quantile predictions (alert when predicted lower bound < 49.8 Hz).

Out-of-Season Validation

Table B.3: December 2019 Performance (Winter Conditions)

Metric	August (Full Month)	December (Testing)	Change	Assessment
PICP (%)	82.1	87.9	+5.8 pp	Improved coverage
Pinball Loss (Lower)	0.00268	0.00251	−6.3%	Robust
MAE Lower (Hz)	0.0208	0.0215	+3.4%	Stable
Calibration ($\alpha=0.1$)	8.9%	5.9%	−3.0 pp	Conservative

Seasonal stability in December indicates the robustness of physics-informed features. The model generalizes well to winter generation/demand patterns despite being trained exclusively on August data.

Computational Performance

Table B.4: Real-Time System Latency

Operation	Latency	Requirement	Status
Data refresh (1 second)	0.12s	<1.0s	Pass
Model inference	0.08s	<0.5s	Pass
SHAP explanation	0.15s	<0.5s	Pass
Dashboard render	0.05s	<1.0s	Pass
Total cycle	0.40s	<2.0s	Pass

Measurements conducted on a Linux-based development workstation (8-core CPU), 16GB RAM, SSD storage. Python 3.11, Polars 0.20.x, LightGBM 4.1.x.

Feature Importance Stability

Table B.5: Top Feature Importance Across Folds

Feature	00:00–08:00	08:00–16:00	16:00–24:00	Mean (%)
Grid Frequency	85.2	88.4	85.7	86.4
Lag 1s	6.1	4.9	7.8	6.3
Wind Speed	2.2	1.7	0.9	1.6
Hour of Day	3.0	1.2	1.6	1.9
Lag 60s	0.7	0.9	0.9	0.8

CV = Coefficient of Variation ($\text{Std}/\text{Mean} \times 100$). Low CV (<5%) for top features indicates stable importance across cross-validation folds.

August 9, 2019 Blackout Event Metrics

Table B.6: Blackout Prediction Timeline

Time (UTC)	Event	Actual Freq (Hz)	Pred Lower Bound (Hz)	Advance Warning (s)	Status
15:52:30	Pre-disturbance	50.006	49.982	—	Stable
15:52:35	Actual threshold breach	49.790	49.970	—	Event
15:52:40	Model alert triggered	49.372	49.796	−5	Warning
15:53:49	Actual Nadir reached	48.800	—	69	Blackout

Alert threshold: 49.80 Hz. Actual nadir: 48.80 Hz. Automatic load shedding triggered at 48.80 Hz.

8.3 Appendix C: Dashboard User Manual and Interface Documentation

This appendix provides complete documentation for the GridGuardian Control Room dashboard, including user instructions and annotated screenshots.

System Overview

The GridGuardian Control Room is a real-time monitoring and predictive analytics dashboard for UK power grid stability. It integrates:

- Live frequency data from NESO CKAN API
- Weather data from Open-Meteo API
- Machine learning predictions (10-second horizon)
- SHAP explainability visualisations
- Intervention simulation capabilities

Target Users: Grid control room operators, energy traders, system planners

Access Requirements: Web browser (Chrome/Firefox/Safari), Python 3.11+ runtime

Navigation Controls

Time Navigation

Control	Function	Default	Range
Date Range Picker	Select overall analysis period	Last 7 days	2019-01-01 to present
Date Picker	Jump to specific day	Today	Within selected range
Step Back/Forward (<< < > >>)	Navigate by 1 second or 1 minute	—	—
Exact Time Input	Enter HH:MM:SS UTC	Current	00:00:00–23:59:59
Autoplay	Automatic time advance	Off	Speed: 1×, 5×, 10×, 30×, 60×

Autoplay mode replays historical events at accelerated speed for training and post-event analysis.

Playback Controls

- **Play/Pause:** Start/stop autoplay
- **Speed Selector:** Adjust playback rate (1× = real-time, 60× = 1 minute per second)
- **Loop Toggle:** Repeat playback when reaching end of selected range
- **Go to Blackout:** Jump directly to August 9, 2019, 15:52:00 UTC (preset for training)

Alert Configuration Panel

Prediction Horizon

Setting	Description	Default	Range
Time to Alert	Displayed warning lead time	10 seconds	5–60 seconds

Note: The underlying model always predicts at the trained 10-second horizon. This slider adjusts the displayed label only, allowing operators to evaluate different warning thresholds.

Instability Threshold

Setting	Description	Default	Range
Frequency Threshold	Alert trigger level	49.80 Hz	49.50–49.90 Hz

Dual-Trigger Logic:

- **Condition A:** Predicted lower bound < threshold (proactive alert)
- **Condition B:** Current frequency < threshold (reactive alert)
- Alert activates when Condition A **OR** Condition B is true


Recommended settings:

- **Conservative:** 49.85 Hz (earlier warning, more false positives)
- **Standard:** 49.80 Hz (balanced)
- **Aggressive:** 49.75 Hz (later warning, fewer false positives)

Main Frequency Display

Visual Elements

Figure C.1: Frequency Plot Components



Legend:

- **Cyan Line:** Actual measured frequency (1-second resolution)
- **Orange Shaded Band:** Prediction uncertainty (10th–90th percentile)
- **Orange Line:** Predicted lower bound (10th percentile)
- **Red Dashed:** Alert threshold (configurable, default 49.80 Hz)
- **Red Dot:** Current time marker
- **Green Zone:** Safe operation (49.80–50.20 Hz)
- **Yellow Zone:** Caution (49.50–49.80 Hz)
- **Red Zone:** Emergency (below 49.50 Hz)

Alert States

State	Visual Indicator	Audio	Meaning
Stable	Green background	None	Frequency within safe bounds, predictions confident
Caution	Yellow background	Single tone	Predicted lower bound approaching threshold
Warning	Orange background	Repeated tone	Alert threshold crossed, intervention recommended
Critical	Red background	Continuous alarm	Current frequency below threshold, emergency response activated
Model Uncertainty	Purple border	—	High prediction uncertainty, trust with caution

SHAP Risk Drivers Panel

Explanation Display

The SHAP panel provides real-time explanations of model predictions using horizontal bar charts.

Figure C.2: SHAP Risk Drivers Example (August 9, 2019)

Risk Drivers (pushing frequency down ↑)

RoCoF (5s smoothed)  -0.042 Hz
└ Rapid frequency decline

Wind Ramp Rate (OpSDA)  -0.031 Hz
└ Sudden generation loss

Renewable Penetration  -0.018 Hz
└ Low inertia vulnerability

Time of Day (15:52)  -0.008 Hz
└ Evening peak demand

Base Value: 50.00 Hz Total: -0.099 Hz
Predicted Lower Bound: 49.901 Hz

Stabilizing Factors (pushing frequency up ↓)

Temperature | +0.003 Hz
└ Reduced heating demand

Interpretation Guidelines:

- **Red bars:** Features increasing instability risk (pushing frequency down)
- **Blue bars:** Features reducing risk (pushing frequency up)
- **Bar length:** Magnitude of contribution
- **Bar direction:** Sign of contribution (left=negative, right=positive)

Operator Interpretation

When an alert triggers, operators should examine:

1. **Dominant drivers:** Which features contribute most to the prediction?
2. **Physical plausibility:** Do the drivers make sense given current conditions?
3. **Actionable insights:** What interventions address the identified drivers?

Example Alert Interpretation: > “Alert triggered due to negative RoCoF (-0.042) indicating rapid frequency decline, combined with OpSDA wind ramp (-0.031) showing sudden generation loss. High renewable penetration (-0.018) indicates low system inertia. Recommended action: activate synthetic inertia reserves.”

Intervention Simulator

Controls

Control	Function	Default	Range
Synthetic Inertia Slider	Adjust injected power	0 MW	0–5,000 MW
Inertia Constant (H)	System inertia parameter	4.0 s	2.0–8.0 s
Reset Button	Clear simulation	—	—

Physical Model

The simulator recalculates predictions using the swing equation:

$$\Delta f = \frac{\Delta P \times f_0}{2 \times H \times S_{\text{base}}}$$

- Where:
- **ΔP** = injected synthetic inertia (MW) [slider value]
 - **f₀** = nominal frequency (50 Hz) [constant]
 - **H** = system inertia constant (s) [configurable]
 - **S_{base}** = total system capacity (35,000 MW) [constant]

Usage Example

Scenario: Alert triggered, predicted lower bound = 49.75 Hz (0.05 Hz below threshold)


Question: How much synthetic inertia is required to restore safety?

- Procedure:**
1. Set H = 4.0 s (current system estimate)
 2. Gradually increase synthetic inertia slider
 3. Observe predicted lower bound rising in real-time
 4. At ΔP = 1,500 MW: predicted lower bound = 49.82 Hz (above threshold)

Conclusion: 1,500 MW synthetic inertia injection would restore predicted stability.

Visualization

Figure C.3: Intervention Simulation Effect



Model Health Metrics Panel

Quantile Metrics

Metric	Current	Target	Status
Pinball Loss (Lower)	0.00268	<0.02	Pass
Pinball Loss (Upper)	0.00260	<0.02	Pass
PICP	82.1%	≥80%	Pass
MPIW	0.0387 Hz	<0.2	Pass

Calibration Metrics

Quantile	Expected	Observed	Deviation	Status
α=0.1	10.0%	8.9%	−1.1 pp	Robust
α=0.9	90.0%	91.0%	+1.0 pp	Robust

Well-calibrated predictions across quantiles demonstrate reliable uncertainty quantification.

Binary Classifier Metrics

Metric	Value	Target	Status
AUC-ROC	0.978	>0.90	Pass

Troubleshooting

Issue	Possible Cause	Resolution
No data displayed	API connection failure	Check internet connection; verify API status
Stale data	Caching issue	Clear browser cache; restart dashboard
Slow performance	Large date range selected	Reduce date range to <7 days
SHAP not loading	Model file missing	Verify model.pkl exists in /models directory
Alerts not triggering	Stale cache or unit mismatch	Clear Parquet processed_cache/ and retrain models

Keyboard Shortcuts

Shortcut	Function
Space	Play/Pause autoplay
← / →	Step backward/forward 1 second
Shift+← / Shift+→	Step backward/forward 1 minute
Home	Jump to start of selected range
End	Jump to end of selected range
B	Jump to August 9, 2019 blackout
R	Reset intervention simulator
H	Toggle help overlay

Note: The R shortcut may conflict with browser page refresh in some browsers; use Ctrl+R if the dashboard does not respond.

8.4 Appendix D: Dashboard Architecture & Cache Design


This appendix details the software engineering architecture of the “Command Deck” v3, focusing on the split-cache pattern required for real-time grid monitoring.

Data Flow and Caching Logic

The system employs a dual-tier caching strategy to manage the trade-off between data freshness and inference latency:

1. Static Resource Cache (`@st.cache_resource`):
 - **Purpose:** Persists heavy model objects across user sessions.
 - **Managed Assets:** LightGBM models, LSTM weights, Isotonic Calibrators, and Scikit-learn Scalers.
 - **Benefit:** Eliminates the 8s disk I/O bottleneck during UI refreshes.
2. Temporal Data Cache (`@st.cache_data`):
 - **Purpose:** Caches processed time-series data while allowing for real-time invalidation.
 - **Managed Assets:** Parquet-backed frequency streams, Weather API responses, and engineered feature sets.
 - **Validation:** Uses source-code hashing to detect changes in the preprocessing logic, ensuring that feature engineering updates are immediately reflected in the UI.

System Architecture Diagram



Cache Invalidation Strategy

To prevent memory leaks and “stale prediction” errors, the system implements a TTL (Time-To-Live) strategy on the temporal cache:

- Frequency Data: 60-second TTL (forced refresh for real-time alignment).
- Weather Data: 1-hour TTL (hourly resolution matching source data).
- SHAP Explainer: Generated per-row to ensure attributions match the current feature state exactly, bypassing the cache for maximum transparency.