

SCHOOL OF ARCHITECTURE, COMPUTING AND ENGINEERING

Department of Engineering and Computing

GridGuardian: Evaluating Physics-Informed Machine Learning for Grid Instability Early-Warning Systems

Fatema Doctor

2604383

A report submitted in part fulfilment of the degree of
BSc (Hons) in *Data Science and Artificial Intelligence*

Supervisor: Ms. Dhara Parekh
CN6000

07 May 2026

Abstract

Declining system inertia from inverter-based renewables alters UK grid dynamics, compressing the window for corrective intervention to sub-second scales. This dissertation evaluates GridGuardian, a prototype for the probabilistic forecasting of grid instability. The system integrates LightGBM quantile regression with physics-based heuristics, including the Swinging Door Algorithm (OpSDA) for wind ramp extraction, inertia proxies, and post-hoc isotonic recalibration. This hybrid architecture addresses the limitations of traditional machine learning in safety-critical environments. Evaluation was conducted using a 1-second telemetry dataset from the National Energy System Operator (NESO), focused on the August 2019 blackout event. While the tree-based architecture achieved sub-second inference latency ($< 0.20s$) and a high recall rate on the hold-out set, forensic stress-testing revealed significant performance ceilings. The model achieved a Prediction Interval Coverage Probability (PICP) of 73.5%, falling short of the nominal 80% safety target, and exhibited reactive rather than predictive behavior during the final frequency nadir. Seasonal analysis further identifies a need for dynamic recalibration to maintain reliability across shifting grid topologies. This work establishes the performance ceiling of physics-informed tree ensembles for sub-second grid forecasting and identifies calibration reliability as the primary barrier to operational deployment.

Acknowledgments

I would like to express my sincere gratitude to my supervisor, Ms. Dhara Parekh, for her invaluable guidance, continuous feedback, and patience throughout the research and development of this dissertation.

I am also grateful to the University of East London and the open-data providers, specifically the National Energy System Operator (NESO) and Open-Meteo, without whose resources this project would not have been possible.

Finally, my deepest appreciation goes to my parents, family and friends. Thank you for your endless encouragement, understanding, and unwavering support during the many long hours spent working on this project. I could not have reached this milestone without you.

Contents

Abstract	i
Acknowledgments	ii
1 Introduction	1
1.1 Background: The Inertia Crisis and Structural Transformation	1
1.2 Problem Statement and Operational Constraints	1
1.3 Research Question and Hypotheses	2
1.4 Research Aims and Technical Objectives	3
1.5 Report Structure	3
2 Literature Review	5
2.1 Introduction	5
2.2 Global Perspectives and Comparative Grid Failures	5
2.2.1 The Texas ERCOT 2021 Winter Storm	5
2.2.2 The South Australian Blackout (2016)	5
2.3 Power System Physics and the Inertia Challenge	6
2.3.1 The Historical Evolution of Frequency Control	6
2.3.2 The Swing Equation and Inverter-Based Resources	6
2.3.3 Empirical Evidence: The August 2019 Blackout	7
2.3.4 Frequency Response Mechanisms	7
2.4 Machine Learning Methodologies for Power Systems	7
2.4.1 Predictive Horizons in Power Systems	7
2.4.2 Sequential Modeling and Deep Learning	8
2.4.3 Topology-Aware Graph Neural Networks (GNNs)	8
2.4.4 Gradient Boosting and Tree-Based Ensembles	8
2.4.5 Probabilistic Forecasting and Conformal Prediction	9
2.5 Feature Engineering and Physics-Informed Machine Learning (PIML)	10
2.5.1 Automated vs. Domain-Specific Feature Extraction	10
2.5.2 Physics-Informed Neural Networks (PINNs)	10
2.5.3 Physics-Informed Feature Engineering	10
2.6 Explainable AI (XAI) in Safety-Critical Infrastructure	11
2.6.1 The Transparency Requirement	11
2.6.2 The Mathematical Axioms of SHAP	11
2.7 Synthesis and Research Gaps	12
3 Methodology	13
3.1 Introduction	13
3.2 Software Development Life Cycle (SDLC)	13
3.3 Data Acquisition and Pre-processing	14
3.3.1 Asynchronous Data Alignment and Leakage Prevention	14
3.3.2 Hardware Portability and Latency Benchmarking	14
3.4 Feature Engineering	15

3.4.1	The Swinging Door Algorithm (OpSDA) Compression	15
3.4.2	Target Variable and Evaluative Horizon	15
3.5	Machine Learning Architecture	15
3.5.1	Temporal Hold-Out Validation and Calibration Splits	15
3.5.2	Post-hoc Isotonic Recalibration	15
3.5.3	Baseline Architecture: LSTM with Monte Carlo Dropout	16
4	Findings	17
4.1	Introduction	17
4.2	Predictive Accuracy and Error Distribution	17
4.2.1	Residual Error and Quantile Metrics Analysis	17
4.2.2	Classification Performance Metrics	18
4.3	Feature Importance: Validating the Physics Integration	19
4.3.1	Structural Dominance	19
4.4	Probabilistic Forecasting Reliability	19
4.4.1	Isotonic Recalibration Performance	19
4.5	Forensic AI (XAI) and The August Blackout Reconstruction	20
4.5.1	The Fragility Fingerprint: Reactive Identification	21
4.5.2	SHAP Waterfall Decomposition	22
5	Evaluation and Discussion	23
5.1	Introduction	23
5.2	Methodological Defense: Target Variable and Predictive Potential	23
5.3	Impact Analysis: Operational Viability and Visualization	23
5.4	Socio-Economic Impacts and Ethical Considerations	24
5.4.1	Protection of Infrastructure and Forensic Diagnostics	24
5.4.2	Ethical Deployment of Automated Intervention	25
5.5	Out-of-Season Performance and Calibration Drift	25
5.6	Global Interpretability and XAI Stability	25
5.6.1	Feature Magnitude and Decision Boundaries	26
5.7	Architectural Justification: Ablation Study	26
5.8	System Limitations	27
5.9	Recommendations for Future Research	27
6	Conclusion	28
6.1	Summary of Technical Contributions and Hypotheses Resolution	28
6.2	Professional Reflection	29
	Bibliography	30
7	Appendices	32
7.1	Appendix A: Initial Project Proposal	32
7.2	Appendix B: Ethical Approval	34
7.3	Appendix C: Dashboard Outputs	39
7.3.1	C.1: Normal Operational State (Pre-Fault)	39

7.3.2	C.2: High Volatility & Model Uncertainty	40
7.3.3	C.3: Critical Instability Alert (August 9, 2019 Nadir)	40
7.3.4	C.4: Post-Fault Recovery Monitoring	41
7.3.5	C.5: Model Health & Calibration Diagnostics	42
7.4	Appendix D: Code Snippets	43
7.4.1	D.1: Physics-Informed Feature Engineering	43
7.4.2	D.2: Probabilistic Model Logic	45
7.4.3	D.3: Dashboard Engine & Multi-Signal Alert Logic	46
7.4.4	D.4: Global Configuration	47

List of Figures

Figure 1	High-level system architecture of the GridGuardian predictive prototype, detailing asynchronous data fusion, feature engineering, and the dual-model inference path.	13
Figure 2	Residual distributions for the LightGBM models. The Mean Absolute Error (MAE) of 0.0243 Hz represents the realistic performance ceiling of the current prototype.	18
Figure 3	Global feature importance for the LightGBM inference engine, highlighting the dominance of physics-based momentum indicators over raw environmental context.	19
Figure 4	Calibration reliability diagrams demonstrating the alignment of forecasted probability quantiles with empirical observation frequencies.	20
Figure 5	Raw frequency time-series of the August 2019 blackout event overlaid with the LightGBM predicted quantile boundaries, illustrating the reactive alert identification and quantile bound behaviour during the collapse.	21
Figure 6	SHAP Waterfall decomposition for the reactive alert at 15:54:04 UTC. The model identifies extreme RoCoF as the primary driver, confirming physical consistency but highlighting the lack of predictive “early-warning” indicators.	22
Figure 7	Hypothetical simulated intervention timeline illustrating the lead-time margin that would be required to trigger EFR battery deployment within the 1.0-second constraint. This figure contextualises the reactive alert produced by the prototype during the August 2019 event against the operational requirement, rather than presenting a measured predictive outcome.	24
Figure 8	Global SHAP beeswarm plot detailing how feature magnitudes (e.g., extreme negative RoCoF) mathematically push the model towards an unstable frequency prediction.	26
Figure 9	Ethical Approval Form (Page 1)	34
Figure 10	Ethical Approval Form (Page 2)	35
Figure 11	Ethical Approval Form (Page 3)	36
Figure 12	Ethical Approval Form (Page 4)	37
Figure 13	Ethical Approval Form (Page 5)	38
Figure 14	Dashboard displaying stable grid conditions prior to the fault event.	39
Figure 15	The system detecting early-warning turbulence, triggering a “High Model Uncertainty” state.	40
Figure 16	Dashboard alert during the August 9, 2019 event. Note: This visualization represents a high-sensitivity simulation mode (captured at 15:52:34 UTC) used for demonstration; the primary operational alert reported in Chapter 4 occurred at 15:54:04 UTC.	40
Figure 17	The recovery phase, demonstrating the model tracking the post-nadir stabilization.	41

Figure 18 The “Model Health” tab providing continuous transparency into the ML engine’s performance.	42
---	----

List of Tables

Table 1	Comparative summary of forecasting horizons and corresponding computational constraints in power systems.	8
Table 2	Comparative evaluation of physics-informed integration strategies for grid forecasting.	11
Table 3	Agile Sprint Timeline and Project Management Risk Register.	14
Table 4	Comprehensive statistical evaluation metrics for the LightGBM quantile bounding models on the August 2019 dataset. †Pre-recalibration Pinball Loss = 0.0032. The post-recalibration value of 0.0012 is reported in Table 7 (Summer column) as part of the seasonal evaluation, which also reflects isotonic-adjusted bounds. ‡Bootstrap confidence intervals (B=1,000 resamples) computed over the full August 2019 test set. PICP bounds computed using the Wilson score interval.	17
Table 5	Binary classification performance metrics for the alert system on the August 9 hold-out test set ($N = 86,400$ records). The reported recall of 99.2% reflects the system’s sensitivity to structural failure within the test period. However, this high sensitivity is partially due to the heuristic alignment between the instability labels and the physical signals used as features, as discussed in Section 5.2.	18
Table 6	Signal state analysis tracing the multi-physics ‘Fragility Fingerprint’ during the August 9, 2019 blackout.	21
Table 7	Comparative performance and distributional shift evaluation when the summer-trained model is applied to winter grid conditions. †Post-recalibration Pinball Loss. The pre-recalibration Pinball Loss baseline (0.0032) is reported in Table 4.	25
Table 8	Ablation study comparing the inference latency and predictive recall of isolated architectures vs. the hybrid system.	26

Chapter 1: Introduction

1.1 Background: The Inertia Crisis and Structural Transformation

The United Kingdom's electricity system is undergoing a significant structural shift. Driven by net-zero emission commitments and the economics of renewable generation, the National Energy System Operator (NESO) is displacing traditional, fossil-fuel-based power stations in favour of inverter-based resources (IBRs) such as wind and solar photovoltaics. The closure of the Ratcliffe-on-Soar power station in September 2024 marked the end of 142 years of coal-powered electricity in Britain, signifying a major change in the nation's energy landscape.

However, this decarbonisation has introduced a complex engineering challenge regarding power system stability. Historically, the stability of the grid's electrical frequency (typically maintained at 50 Hz) was guaranteed by the physical properties of massive synchronous generators. The turbines of coal, gas, and nuclear plants contain immense rotating mass. This stored kinetic energy acts as a mechanical shock absorber, resisting sudden deviations in frequency whenever there is an imbalance between electricity supply and demand. This phenomenon is known as "synchronous inertia". Under this traditional paradigm, frequency deviations following a power station failure were relatively slow. Grid operators had a buffer—often measured in minutes—to detect the failure and intervene.

Inverter-based renewable sources completely lack this mechanical inertia. They connect to the transmission network via power electronics, not through direct, massive electromechanical coupling. As conventional generation continues to be retired, the aggregate physical momentum of the entire UK grid is plummeting. Data from the National Energy System Operator indicates that system inertia has declined by over 50% since 2010 (National Grid ESO, 2020), and projections indicate it will routinely fall below the critical threshold of 100 GVA-s during low-demand, high-wind periods.

This phenomenon alters the mathematical dynamics of the grid. Governed by the physical "Swing Equation," a grid with lower inertia experiences steeper frequency trajectories during a fault. A disturbance that would have caused a manageable frequency sag twenty years ago now causes a rapid frequency drop. The Rate of Change of Frequency (RoCoF) increases, reducing the time available for corrective action from minutes down to seconds. If the grid frequency drops below the operational boundary of 48.8 Hz, automated protection relays will execute "under-frequency load shedding," blacking out regions of the country to save the physical infrastructure from severe damage.

1.2 Problem Statement and Operational Constraints

The engineering reality of the low-inertia grid is that traditional, reactive monitoring systems are no longer viable. A reactive system waits until a hard frequency threshold (e.g., 49.8 Hz) is breached before sounding an alarm or triggering a response. In a grid with high inertia, a

breach of 49.8 Hz provided ample time to deploy reserve power before the system collapsed to 48.8 Hz. In a low-inertia grid, the collapse from 49.8 Hz to 48.8 Hz can occur in a matter of seconds, rendering reactive alerts functionally useless.

This vulnerability was exposed during the August 9, 2019 UK blackout. Following a lightning strike, a cascading failure disconnected 1,481 MW of generation. The system frequency fell to 48.8 Hz in under 10 seconds, triggering automated load shedding that affected 1.1 million customers, paralyzed national railway networks, and caused backup generator failures at critical infrastructure facilities. The technical post-mortem identified that the grid’s protection mechanisms were calibrated for historical inertia margins that were no longer valid. The system momentum was insufficient to arrest the frequency decay prior to the load shedding. To combat this, NESO procures automated, rapid-response stability services. While traditional Firm Frequency Response (FFR) requires up to 10 seconds to inject its maximum active power, modern Enhanced Frequency Response (EFR) batteries are mandated to achieve 100% active power injection within exactly 1.0 seconds of an automated trigger.

The core engineering challenge is to investigate whether a predictive alerting system can forecast frequency instability before nominal thresholds are breached. For such a system to be operationally viable, the forecast should ideally provide a lead time that exceeds the 1.0-second mechanical deployment constraint of EFR battery systems. This research evaluates the GridGuardian prototype against these benchmarks, identifying the trade-offs between predictive lead-time, interval reliability (PICP), and the transparency requirements of safety-critical automated interventions.

1.3 Research Question and Hypotheses

This research explicitly investigates the following primary question: “To what extent can a physics-informed, tree-based machine learning architecture provide a sub-second, probabilistic forecast of power grid instability that satisfies the lead time and reliability requirements of modern Enhanced Frequency Response systems?”

From this question, the following hypotheses are formulated and rigorously evaluated:

Primary Hypothesis (H_1): By engineering features derived from power system physical heuristics (such as RoCoF and inertia proxies), a tree-based quantile regression model can be configured to target the lower boundary of grid frequency 10 seconds into the future. It is hypothesised that this architecture can maintain a Prediction Interval Coverage Probability (PICP) $\geq 80\%$ under in-distribution conditions (evaluated on the August 2019 hold-out dataset as the primary benchmark) and provide an empirical alert margin of ≥ 1.0 second prior to failure during discontinuous events.

Null Hypothesis (H_0): The integration of physical heuristics into a quantile regression model yields no statistically significant improvement in predictive lead time or algorithmic recall compared to a deep recurrent baseline (LSTM), or the computational and calibration

requirements of the pipeline render it unable to consistently meet industrial safety targets across diverse grid conditions.

1.4 Research Aims and Technical Objectives

The aim of this research is to develop and critically evaluate “GridGuardian”—a functional, physics-informed software prototype for predictive grid monitoring. The project seeks to determine whether a lightweight, interpretable ensemble architecture can bridge the gap between reactive physical thresholds and the predictive requirements of low-inertia grid operations.

The execution of this aim requires the fulfillment of five specific, highly technical objectives:

1. **High-Velocity Data Engineering:** Construct a robust, sub-second data ingestion pipeline capable of merging asynchronous data streams (1-second NESO telemetry and hourly Open-Meteo meteorological data). Implement a strictly backward-filling temporal alignment algorithm (`join_asof`) to mathematically guarantee the prevention of forward-looking data leakage during model training.
2. **Physics-Informed Feature Architecture:** Reject the automated extraction of arbitrary statistical artifacts. Instead, engineer features strictly aligned with grid mechanical realities, specifically implementing discrete digital filtering for RoCoF calculation, the Swinging Door Algorithm (OpSDA) for wind ramp rate extraction, and explicit metrics derived from half-hourly system inertia data.
3. **Probabilistic Inference Configuration:** Deploy a LightGBM gradient boosting engine optimized for quantile regression via the asymmetric Pinball Loss function, targeting the generation of strict 10th and 90th percentile probability bounds to quantify forecast uncertainty.
4. **Algorithmic Transparency Audit:** Implement TreeSHAP (SHapley Additive exPlanations) to audit the internal logic of the model. Evaluate the output for logical consistency—specifically investigating the feature attribution direction and stability across the fault window—to determine its suitability for high-stakes decision support.
5. **Empirical Forensic Stress-Testing:** Subject the prototype to a forensic reconstruction of the August 9, 2019 blackout. Evaluate the realized predictive margin and interval reliability against the strict operational constraints of the UK grid to identify the system’s current performance ceiling. Conduct a formal ablation study to mathematically justify the latency vs. recall trade-offs inherent in the chosen tree-based architecture against deeper recurrent baseline models.

1.5 Report Structure

The remainder of this dissertation documents the engineering, evaluation, and critique of the GridGuardian system. Chapter 2 presents a literature review, critiquing methodological approaches to transient stability and exposing the limitations of deep neural architectures in sub-second inference environments. Chapter 3 details the methodology, providing mathematical documentation for the data pipeline, the feature extraction logic, and the LightGBM

inference configuration. Chapter 4 presents the findings of the system evaluation, detailing the predictive accuracy and interval reliability across the August 2019 test set. This chapter provides an account of the model's performance, including areas where the prototype meets operational targets and where it falls short of safety targets. Chapter 5 evaluates these findings, analyzing the operational viability of the predictive margin against known computational latencies, defending the model architecture, and outlining system limitations. Finally, Chapter 6 concludes the research, reflecting on the engineering implications and future work.

Chapter 2: Literature Review

2.1 Introduction

The structural transition of electricity grids toward inverter-based renewable energy sources presents a fundamental challenge to frequency stability. The displacement of synchronous generation reduces system inertia, modifying the physical dynamics of frequency response and reducing the time available for operator intervention. This review synthesises current research concerning grid stability in low-inertia environments. It evaluates the comparative efficacy of machine learning forecasting architectures and the integration of physical constraints into data-driven models. The chapter examines the specific operational requirements for predictive alerting systems and the necessity of algorithmic interpretability in safety-critical grid management.

2.2 Global Perspectives and Comparative Grid Failures

While this research focuses on the UK National Grid, the challenges of low-inertia stability are global. Examining international failure modes provides critical context for the necessity of predictive AI systems.

2.2.1 The Texas ERCOT 2021 Winter Storm

A clear example of cascading grid failure occurred in February 2021 in the ERCOT (Electric Reliability Council of Texas) interconnection. The Texas failure represents a distinct failure mode compared to the UK 2019 event: rather than a rapid, transient frequency collapse caused by a singular lightning strike, it was a sustained, systemic failure triggered by extreme cold weather and un-winterized thermal generation assets.

Despite these differences, both events share the same underlying physical vulnerability: a sudden mismatch between load and generation that precipitates a dangerous frequency nadir. During the Texas storm, the grid frequency dropped below the 59.4 Hz operational limit in under four seconds, forcing operators into manual load-shedding to avert a total “black start” scenario. This contrast demonstrates that while generation-load mismatches are universal, their temporal dynamics vary. Consequently, any predictive alerting system must be calibrated to the specific failure modes—whether transient or sustained—of its target network, rather than relying on a generalized detection logic.

2.2.2 The South Australian Blackout (2016)

The South Australian blackout of 2016 represents an early warning for low-inertia grids. Triggered by a severe storm that damaged transmission lines, the grid—which had over 40% wind penetration at the time—experienced a rapid RoCoF that exceeded the capabilities of existing protection relays. The Australian Energy Market Operator (AEMO) (AEMO, 2017) noted that the speed of the frequency collapse was so great that human dispatchers could not respond in time.

This event validates the core premise of GridGuardian: in modern, high-renewable grids, the “decision window” for human operators has effectively closed. The South Australian event demonstrated that the transition from 3.0 Hz/s to 1.0 Hz frequency nadirs occurs in less than 1.5 seconds—a timescale that the GridGuardian prototype did not achieve under the August 2019 conditions, motivating the identification of pre-fault feature architectures as a direction for future research. By comparing the UK context with Australia and Texas, it is clear that the problem of ‘transient fragility’ is a property of the global energy transition, necessitating further development of AI-driven early-warning systems.

2.3 Power System Physics and the Inertia Challenge

2.3.1 The Historical Evolution of Frequency Control

Historically, power system frequency stability was guaranteed by the inherent physical properties of massive synchronous generators (coal, gas, and nuclear plants). The kinetic energy stored in the rotating turbines provided a natural, instantaneous buffer against supply-demand imbalances, a phenomenon formally known as “synchronous inertia.” Under this traditional paradigm, frequency deviations were relatively slow, allowing grid operators to rely on “droop control”—a primary frequency response mechanism where turbine governors mechanically adjusted steam valves in proportion to frequency deviations. The response times required were on the order of tens of seconds, making frequency control a manageable, predominantly mechanical problem.

However, the rapid decarbonisation of the energy sector has rendered this paradigm obsolete. As modern grids transition to renewable energy sources, synchronous mass is actively being displaced.

2.3.2 The Swing Equation and Inverter-Based Resources

Power system frequency dynamics are governed by the swing equation, which mathematically formalises the relationship between kinetic energy in rotating masses and active power imbalances (Kundur et al., 1994). The fundamental equation dictates that the rate of change of frequency ($\frac{df}{dt}$) following a disturbance is inversely proportional to the system’s synchronous inertia (H). Therefore, reductions in inertia directly produce steeper frequency trajectories during generation losses.

(Milano et al., 2018) established that modern networks exhibit increasingly volatile frequency deviations as decentralised, Inverter-Based Resources (IBRs) replace centralised thermal plants. Because IBRs (such as wind and solar photovoltaics) connect to the grid via power electronics rather than direct electromechanical coupling, they inherently provide zero natural synchronous inertia. (Tielens & Van Hertem, 2016) demonstrated that operating below critical inertia thresholds produces Rate of Change of Frequency (RoCoF) values exceeding standard protection limits. This physical constraint dictates that contemporary grid monitoring systems must evaluate continuous frequency derivatives rather than static frequency thresholds alone.

2.3.3 Empirical Evidence: The August 2019 Blackout

The August 9, 2019 UK blackout provides an empirical case study for low-inertia vulnerability. The simultaneous disconnection of the Little Barford gas plant and the Hornsea One wind farm removed 1,481 MW of generation, causing the grid frequency to decline to 48.8 Hz within 10 seconds (Homan et al., 2020).

The subsequent technical investigation by the Office of Gas and Electricity Markets (Ofgem and BEIS, 2019) identified a misalignment between existing protection configurations and actual system inertia. Generators were calibrated for historical inertia margins that were mathematically invalid for a wind-dominated grid profile. This event demonstrated the inadequacy of reactive, threshold-based protection mechanisms when system momentum is insufficient to arrest frequency decay prior to automated load shedding.

2.3.4 Frequency Response Mechanisms

To mitigate declining physical inertia, grid operators must procure synthetic stability services. The National Energy System Operator's operability strategy (National Grid ESO, 2020) categorises these interventions strictly by response velocity. Traditional Firm Frequency Response (FFR) requires up to 10 seconds to reach maximum active power output. In contrast, Enhanced Frequency Response (EFR), typically provided by grid-scale Battery Energy Storage Systems (BESS), must deliver full output within 1 second. The operational utility of any predictive alerting system is strictly defined by these mechanical deployment constraints; a forecasting horizon must perfectly align with the activation parameters of the available response technology.

2.4 Machine Learning Methodologies for Power Systems

As the physical complexity of the grid outpaces the capabilities of traditional analytical solvers, data-driven approaches have gained prominence for dynamic security assessment (Zhao et al., 2019).

2.4.1 Predictive Horizons in Power Systems

A critical distinction in power system forecasting literature is the temporal predictive horizon. The vast majority of existing ML applications focus on "day-ahead" or "hour-ahead" unit commitment forecasting, utilizing autoregressive integrated moving average (ARIMA) models or deep learning to predict overall load curves for energy market trading. These models operate in minutes or hours and prioritize average long-term accuracy.

Conversely, "sub-second" transient stability forecasting addresses a fundamentally different problem: detecting structural collapses instantaneously. Models operating in this domain face extreme latency constraints and must prioritize worst-case boundary predictions over average expected values. The literature reveals a distinct gap where advanced ML techniques successfully deployed for day-ahead forecasting fail dramatically when constrained by sub-second latency budgets.

Horizon	Latency Budget	Typical Architectures	Primary Application
Day-Ahead	Hours	ARIMA, Transformers	Unit Commitment & Trading
Intra-Hour	Minutes	Deep Learning (LSTM)	Economic Dispatch
Sub-Second	< 1 Second	Tree Ensembles (Light-GBM)	Transient Stability & EFR

Table 1: Comparative summary of forecasting horizons and corresponding computational constraints in power systems.

2.4.2 Sequential Modeling and Deep Learning

Recurrent architectures, specifically Long Short-Term Memory (LSTM) networks, have historically dominated sequential forecasting in energy systems due to their capacity to capture complex temporal dependencies (Goodfellow et al., 2016). However, LSTMs suffer from sequential bottlenecks, preventing true parallelization during training and inference.

More recently, Transformer-based architectures have emerged for multi-horizon energy forecasting. The Informer architecture (Zhou et al., 2021) utilizes ProbSparse self-attention to reduce the $O(L^2)$ time complexity of standard transformers, making it effective for long-sequence load forecasting. Similarly, the Temporal Fusion Transformer (TFT) (Lim et al., 2021) integrates specialized interpretable attention heads that align well with grid monitoring requirements. However, while Transformers excel at capturing long-term seasonal cycles, the parameter count and complex attention mechanisms introduce high inference latency when attempting to process 1-second telemetry for instantaneous transient stability alerting.

2.4.3 Topology-Aware Graph Neural Networks (GNNs)

Modern power grids are not merely isolated time-series streams; they are highly interconnected spatial graphs. Graph Neural Networks (GNNs) have gained significant traction for topology-aware dynamic security assessment (Liao et al., 2021). By representing buses as nodes and transmission lines as edges, algorithms like Graph Convolutional Networks (GCN) can mathematically model how a localized fault (e.g., a lightning strike on a specific line) propagates spatially across the network. While GNNs offer a more holistic representation of grid physics than purely temporal models, they require complete, real-time observability of the entire grid topology—data that is rarely available at sub-second resolution outside of strict academic simulations. For edge-deployed alerting systems constrained to localized frequency telemetry, temporal ensembles remain the more viable architecture.

2.4.4 Gradient Boosting and Tree-Based Ensembles

While deep neural networks suffer from high inference latency, tree-based ensemble methods provide a computationally efficient alternative for tabular forecasting. However, not all tree architectures are suitable for sub-second control room deployment.

Random Forest (RF) constructs multiple decision trees independently using bootstrap aggregation (bagging). While highly robust against overfitting, RF requires deep trees to capture

complex non-linearities. During inference, traversing hundreds of deep, independent trees simultaneously introduces unacceptable computational bottlenecks. Furthermore, standard RF cannot natively optimize for asymmetric loss functions like quantile bounds without significant post-hoc structural modifications.

Extreme Gradient Boosting (XGBoost) resolves the depth issue by building trees sequentially, minimizing residual errors. However, XGBoost utilizes a level-wise (depth-first) tree growth strategy. When processing the 2.6 million high-resolution data points required for continuous grid monitoring, XGBoost's exact greedy algorithm for finding optimal split points becomes memory-bound.

The LightGBM algorithm (Ke et al., 2017) resolves these specific bottlenecks, making it uniquely suited for the strict 1-second inference constraints of EFR battery deployment. LightGBM utilizes histogram-based decision tree construction, binning continuous variables into discrete buckets. This reduces memory allocation by over 80% compared to exact greedy algorithms. Critically, LightGBM employs a leaf-wise (best-first) growth strategy, expanding the leaf with the maximum delta loss rather than growing level-by-level. This produces deeper, more asymmetrical trees that achieve lower error rates with significantly fewer splits. Furthermore, Gradient-based One-Side Sampling (GOSS) allows LightGBM to selectively drop data instances with small gradients during training, maintaining predictive accuracy while drastically reducing the computational burden.

(Qiu et al., 2020) demonstrated the efficacy of ensemble machine learning for frequency response prediction, noting that tree-based ensembles excel when provided with carefully engineered, domain-specific features. Because the histogram-based tree traversal of LightGBM is computationally trivial compared to backpropagation networks or deep Random Forests, it emerges as the optimal architecture for real-time edge deployment.

2.4.5 Probabilistic Forecasting and Conformal Prediction

Operational grid management necessitates uncertainty quantification. Grid dispatchers must evaluate the probability of a fault against the economic cost of intervention. Conventional implementations of gradient boosting default to point-estimation (minimizing Mean Squared Error), predicting expected values without corresponding probability distributions.

Quantile regression (Koenker & Bassett Jr, 1978) resolves this by computing conditional quantiles, yielding explicit probability intervals. By minimizing the asymmetric Pinball Loss function, models penalize underestimations and overestimations differently (Hastie et al., 2009).

An emerging alternative for uncertainty quantification in safety-critical systems is Conformal Prediction (Vovk et al., 2005). Unlike quantile regression—which requires the model to learn the conditional distribution directly and can suffer from miscalibration—Conformal Prediction is a post-hoc framework that guarantees strict marginal coverage probabilities regardless of the underlying model architecture or data distribution. While mathematically

robust, Conformal Prediction typically generates wider interval bounds than optimized quantile regression, which can increase the false-positive alerting rate in tightly constrained frequency environments. Consequently, the integration of quantile regression into efficient tree-based algorithms currently presents the optimal mechanism for generating real-time, tightly bounded stability forecasts.

2.5 Feature Engineering and Physics-Informed Machine Learning (PIML)

2.5.1 Automated vs. Domain-Specific Feature Extraction

In modern data science, automated feature extraction libraries, such as `tsfresh` (Christ et al., 2018), have gained immense popularity. These algorithms automatically generate thousands of statistical features (e.g., Fourier coefficients, wavelet transforms, kurtosis) from raw time-series data, allowing models to detect hidden patterns without requiring domain expertise.

While statistically powerful, automated extraction is entirely detached from physical reality. A grid operator cannot act upon an alert triggered by a “shift in the third-order wavelet coefficient.” In contrast, domain-specific feature engineering deliberately constructs variables that represent known physical phenomena. Transforming raw frequency into Rate of Change of Frequency (RoCoF) directly encodes the grid’s acceleration—a metric instantly actionable and legally codified in grid operation standards.

2.5.2 Physics-Informed Neural Networks (PINNs)

Standard machine learning models treat power grids as pure statistical distributions, ignoring fundamental mechanical realities. Physics-Informed Machine Learning (PIML) addresses this limitation by embedding domain principles into the learning architecture (Karniadakis et al., 2021).

Physics-Informed Neural Networks explicitly penalise deviations from differential equations (such as the swing equation) during the loss calculation (Raissi et al., 2019). By adding a physical residual term to the objective function, PINNs ensure that predictions obey the laws of conservation. While PINNs demonstrate excellent generalisation in transient stability assessments, enforcing continuous mathematical constraints during real-time, sub-second inference introduces substantial computational overhead.

2.5.3 Physics-Informed Feature Engineering

An alternative, highly efficient strategy involves physics-informed feature engineering. Rather than constraining the model’s internal architecture, the input domain is structurally modified to include explicit physical derivatives. Proxying inertia via renewable penetration ratios allows the model to learn the swing equation indirectly. This approach leverages the sub-second computational efficiency of gradient boosting (Ke et al., 2017) while securely grounding the statistical predictions in mechanical reality.

Methodology	Primary Advantage	Operational Drawback
Pure Data-Driven (Black Box)	High capacity for non-linear mapping	Ignores physical laws; prone to hallucination
Physics-Informed Neural Networks (PINNs)	Mathematically guarantees physical conservation limits	High inference latency; struggles with sub-second bounds
Physics-Informed Feature Engineering	Sub-second inference; utilizes high-speed tree traversal	Requires expert domain knowledge for feature creation

Table 2: Comparative evaluation of physics-informed integration strategies for grid forecasting.

2.6 Explainable AI (XAI) in Safety-Critical Infrastructure

2.6.1 The Transparency Requirement

Black-box models face adoption barriers in power system operations due to the inability to verify the logic preceding automated interventions. Algorithmic transparency is a requirement for integrating machine learning systems into live grid control rooms (Machlev et al., 2022).

2.6.2 The Mathematical Axioms of SHAP

SHAP (SHapley Additive exPlanations) resolves the mathematical inconsistencies of traditional heuristic attribution methods by employing cooperative game theory to distribute exact feature contributions, drawing upon foundational mathematics established by Shapley (Shapley, 1953) and synthesized for machine learning by (Lundberg & Lee, 2017) and (Covert et al., 2021). In safety-critical contexts, an attribution algorithm must be mathematically provable. SHAP is the only additive feature attribution method that mathematically guarantees four critical axioms:

1. **Efficiency:** The feature attributions must sum precisely to the difference between the model's current prediction and the expected baseline prediction. No fractional attribution is lost.
2. **Symmetry:** If two features contribute equally to all possible coalitions, their SHAP values must be identical.
3. **Dummy (Null Effect):** If a feature never changes the predicted value regardless of the coalition it joins, its SHAP attribution is guaranteed to be exactly zero.
4. **Linearity (Additivity across models):** If a model is a linear combination of multiple models (such as an ensemble of trees), the SHAP value of a feature is the linear combination of its SHAP values in the individual models.

These mathematical guarantees distinguish SHAP from heuristic explanation methods like LIME, which approximate local decision boundaries and can suffer from severe local instability. Furthermore, exact tree-based SHAP implementations (TreeSHAP) leverage the internal split structures of tree ensembles to compute these values in low-order polynomial time,

providing distinct operational latency advantages over the slow, model-agnostic permutation methods required by neural networks.

However, modern XAI evaluation demands an analysis of robustness and adversarial vulnerability. (Slack et al., 2020) demonstrated that post-hoc explanation methods like SHAP can be manipulated or become mathematically unstable when presented with highly correlated out-of-distribution data. When dealing with highly correlated grid features (e.g., wind speed and renewable generation percentage), SHAP values can exhibit temporal instability, distributing attribution arbitrarily between correlated variables. If an attribution algorithm requires multiple seconds to compute, or if its explanations oscillate wildly during a cascading fault, its utility to a grid operator is zero.

2.7 Synthesis and Research Gaps

This critical review identifies three persistent gaps in contemporary grid stability literature:

1. **Absence of Real-Time Probabilistic Bounds:** While deep learning architectures offer strong point-prediction accuracy for day-ahead markets, the sub-second generation of reliable, calibrated quantile bounds remains limited, restricting risk-aware decision making in the control room during transient instability.
2. **Computational Feasibility of Physics Constraints:** PINNs successfully embed physical laws but struggle to meet the strict inference constraints of EFR battery deployment. The efficacy of physics-informed feature engineering coupled with high-speed tree ensembles remains inadequately benchmarked against formal neural architectures.
3. **Adversarial and Discontinuous Event Validation:** The majority of forecasting models and their XAI explanations are evaluated on normal, continuous grid operations. Systematic validation of model predictions and SHAP stability against severe, discontinuous faults (such as the August 2019 blackout) is rare.

The present research addresses these specific methodological deficiencies by engineering GridGuardian: a hybrid, physics-informed LightGBM architecture designed to generate probabilistic frequency boundaries, explicitly validated against the non-linear dynamics of a major historical blackout.

Chapter 3: Methodology

3.1 Introduction

This chapter details the methodology implemented to engineer the GridGuardian alerting system. The research follows an empirical software engineering approach, translating established power system physics into a functional machine learning prototype. The methodology is constrained entirely to the techniques, architectures, and data processing routines actively deployed in the project codebase. The complete implementation, including all data ingestion scripts, feature engineering pipelines, trained model artefacts, and dashboard outputs, is publicly available at: <https://github.com/Fatema-Dr/UK-Grid-instability>. This codebase serves as the definitive evidence of the project's technical execution, verifying that GridGuardian is a functional engineering prototype rather than a conceptual model.




Figure 1: High-level system architecture of the GridGuardian predictive prototype, detailing asynchronous data fusion, feature engineering, and the dual-model inference path.

3.2 Software Development Life Cycle (SDLC)

An Agile methodology managed the project lifecycle to accommodate the iterative nature of machine learning development. The project was structured into three distinct two-week sprints:

Sprint Phase	Key Deliverables	Risk Mitigation
Sprint 1: Data Eng (Weeks 1-2)	NESO/Open-Meteo APIs, Polars join_asof pipeline	Mitigated memory limits via Rust-based Polars over Pandas
Sprint 2: ML Proto-typing (Weeks 3-4)	LightGBM baseline, quantile calibration, LSTM dual-architecture	Mitigated structural over-fitting via strict L1/L2 regularization
Sprint 3: System Integration (Weeks 5-6)	SHAP interpretability module, final threshold logic, empirical validation	Mitigated inference latency via histogram binning in LightGBM

Table 3: Agile Sprint Timeline and Project Management Risk Register.

Version control (Git) was maintained throughout. To prevent look-ahead bias during data aggregation, continuous unit testing verified the chronological integrity of all temporal joins.

3.3 Data Acquisition and Pre-processing

The system processes multi-resolution data from the National Energy System Operator (NESO) and the Open-Meteo API.

3.3.1 Asynchronous Data Alignment and Leakage Prevention

The core of the data fusion pipeline is the Rust-based Polars `join_asof` algorithm. This implementation achieves an alignment latency of less than 15ms for the entire August dataset, a critical technical requirement for sub-second operations. By enforcing a strict backward-filling causal constraint ($t_w \leq t_f$), the system mathematically prevents the ‘forward-looking’ data leakage that often invalidates time-series forecasting projects. This alignment is not a conceptual design but is fully realized in the project’s `src/data_ingestion.py` module.

3.3.2 Hardware Portability and Latency Benchmarking

The system was developed on a Linux-based architecture (Ubuntu 22.04 LTS) to ensure maximum compatibility with grid-scale edge hardware. LightGBM (Ke et al., 2017) was selected as the primary inference engine due to its histogram-based binning and leaf-wise (best-first) growth strategy, which allows for deeper, more asymmetrical trees than traditional depth-first algorithms.

To ensure the ablation study results are operationally relevant, all inference latency benchmarks were conducted strictly on CPU-only hardware (AMD Ryzen 9 5950X), mirroring the power-constrained environment of grid substations where dedicated GPU resources are typically unavailable. The resulting 0.20s inference latency confirms the system’s computational viability for 1Hz telemetry processing, providing a reproducible CPU-only benchmark that approximates the computational constraint of power-limited edge deployments, though formal validation on certified substation hardware remains outside the scope of this prototype.

3.4 Feature Engineering

Feature engineering translates raw telemetry into physically meaningful variables representing grid momentum and stress.

3.4.1 The Swinging Door Algorithm (OpSDA) Compression

To computationally extract the structural wind ramp rate, the system applies the Swinging Door Algorithm (OpSDA) (Bristol, 1990). By converting raw, high-frequency wind velocity into discrete OpSDA ramp rates, the feature engineering pipeline isolates sustained physical momentum shifts from irrelevant localized turbulence, directly providing the LightGBM engine with a clean derivative of meteorological stress.

3.4.2 Target Variable and Evaluative Horizon

A critical distinction must be made between the model's objective function and the system's operational alert output. The LightGBM architecture performs **continuous quantile regression** to probabilistically forecast the 10th and 90th percentile frequency boundaries targeting a 10-second temporal horizon ($t + 10$).

This specific window was selected to test whether a predictive alert could comfortably exceed the 1.0-second mechanical deployment constraint of Enhanced Frequency Response (EFR) batteries. It must be noted that this $t + 10$ horizon is a **design target**; the realized predictive lead-time during a cascading fault is a function of the model's realized sensitivity and the non-linear dynamics of the specific event, and is not a guaranteed constant of the architecture. A future state is flagged with a binary `target_is_unstable` alert if the model's predicted lower quantile breaches any of the following physical failure bounds:

1. **Operational Breach:** The predicted 10th percentile frequency drops below 49.85 Hz.
2. **Momentum Failure:** The predicted 10th percentile frequency drops below 49.95 Hz while the current smoothed RoCoF is worse than -0.02 Hz/s.

3.5 Machine Learning Architecture

3.5.1 Temporal Hold-Out Validation and Calibration Splits

Standard k -fold cross-validation is invalid for time-series data. The evaluation utilizes a strict chronological three-way split to ensure out-of-sample integrity:

1. **Training Set:** August 1–6, 2019 (518,400 records).
2. **Calibration Set:** August 7–8, 2019 (Used for post-hoc isotonic recalibration).
3. **Testing Set:** August 9, 2019 (Including the discontinuous blackout event).

While the training data is restricted to the period before August 9, the author acknowledges that the calibration set (August 7–8) was used for model tuning, representing a risk of **calibration leakage** that is addressed in Chapter 5.

3.5.2 Post-hoc Isotonic Recalibration

Initial evaluation demonstrated that the LightGBM quantile outputs exhibited a pessimistic calibration bias. To rectify this, Post-hoc Isotonic Recalibration (Kuleshov et al., 2018) was

applied. Isotonic regression fits a monotonically increasing mapping function to ensure that nominal quantile levels correspond to true empirical frequencies. This calibration is a post-processing step that does not modify the underlying LightGBM weights, allowing for a transparent account of the model’s ‘raw’ vs. ‘adjusted’ performance ceilings.

3.5.3 Baseline Architecture: LSTM with Monte Carlo Dropout

To establish a performance benchmark, a recurrent Long Short-Term Memory (LSTM) network was implemented. The LSTM utilizes Monte Carlo Dropout (Gal & Ghahramani, 2016) during inference. By maintaining dropout layers in an active state during prediction across 100 stochastic passes, the baseline generates a predictive distribution for comparison. The LSTM baseline was not subjected to extensive hyperparameter optimization; a more tuned LSTM architecture might narrow the performance gap with the tree-based model.

GridGuardian: Scope of Contributions

- **Implemented Components:** Causal Polars data pipeline; Physics-informed feature engineering (OpSDA, RoCoF); LightGBM Quantile Regression; Isotonic Recalibration logic; Proof-of-concept Streamlit Dashboard.
- **Evaluated Metrics:** Statistical accuracy (MAE, Pinball Loss); Interval reliability (PICP); Computational latency (CPU-benchmarked); SHAP stability.
- **Illustrative/Simulated Results:** The “EFR battery intervention timeline” and hypothetical lead-times are derived from simulated reconstructions of the 2019 blackout event based on the model’s reactive and predictive alert triggers.

Chapter 4: Findings

4.1 Introduction

This chapter presents the objective, empirical results generated by the GridGuardian forecasting and alerting system. The evaluation focuses exclusively on the predictive outputs and explainability metrics derived directly from the LightGBM inference engine and the LSTM Monte Carlo comparator.

The August 9, 2019 UK blackout serves as the primary evaluation scenario. The analysis utilizes a strict hold-out testing set, ensuring that all findings presented herein reflect true out-of-sample generalization.

4.2 Predictive Accuracy and Error Distribution

Before evaluating the alerting logic, the fundamental accuracy of the underlying quantile regression models must be quantified. The system is required to capture the structural trajectory of the grid frequency; however, statistical evaluation reveals that the “tightness” of the predicted bounds comes at the expense of nominal coverage reliability.

4.2.1 Residual Error and Quantile Metrics Analysis

To evaluate the LightGBM models (predicting the $\alpha = 0.1$ lower bound), multiple rigorous statistical metrics were calculated. Table 4 presents the primary benchmark metrics for the August 2019 dataset.

Metric	Lower Bound ($\alpha = 0.1$)	Upper Bound ($\alpha = 0.9$)	Operational Target	Status
Pinball Loss [†]	0.0032 $\pm$ 0.0001	0.0160 $\pm$ 0.0004	< 0.02	Pass
Mean Absolute Error (Hz)	0.0243 [0.0231, 0.0255]	0.0193 [0.0183, 0.0203]	< 0.05	Pass
Root Mean Square Error (Hz)	0.0419 [0.0398, 0.0440]	0.0409 [0.0389, 0.0429]	< 0.10	Pass
PICP (%)	73.5 [73.4, 73.6]	—	$\geq$ 80%	Fail
MPIW (Hz)	0.0357 [0.0340, 0.0375]	—	< 0.20	Pass

Table 4: Comprehensive statistical evaluation metrics for the LightGBM quantile bounding models on the August 2019 dataset. [†]Pre-recalibration Pinball Loss = 0.0032. The post-recalibration value of 0.0012 is reported in Table 7 (Summer column) as part of the seasonal evaluation, which also reflects isotonic-adjusted bounds. [‡]Bootstrap confidence intervals (B=1,000 resamples) computed over the full August 2019 test set. PICP bounds computed using the Wilson score interval.

The Mean Absolute Error (MAE) of 0.0243 Hz indicates that the model follows the general grid frequency within a 25mHz margin. However, the Prediction Interval Coverage Probability (PICP) achieved only 73.5%. This indicates that the model’s physical boundary predictions are statistically overconfident, failing the strict 80% coverage criterion. This result suggests that while the tree-based architecture is computationally efficient, it exhibits an “interval collapse” during periods of high volatility, where the predicted bounds fail to encapsulate the true frequency nadir with sufficient regularity for safety-critical deployment.




Figure 2: Residual distributions for the LightGBM models. The Mean Absolute Error (MAE) of 0.0243 Hz represents the realistic performance ceiling of the current prototype.

4.2.2 Classification Performance Metrics

Because the continuous quantile bounds are heuristically converted into a binary `target_is_unstable` alert, classification metrics can be evaluated on the hold-out set containing the August 2019 blackout (Table 5).

Metric	Value	95% Confidence Interval
Recall (Sensitivity)	99.2%	[98.5%, 99.7%]
Precision	94.1%	[92.8%, 95.3%]
F1-Score	96.6%	[95.6%, 97.4%]
ROC-AUC	0.998	[0.996, 0.999]

Table 5: Binary classification performance metrics for the alert system on the August 9 hold-out test set ($N = 86,400$ records). The reported recall of 99.2% reflects the system’s sensitivity to structural failure within the test period. However, this high sensitivity is partially due to the heuristic alignment between the instability labels and the physical signals used as features, as discussed in Section 5.2.

The system achieves a 99.2% recall rate on the test set. This confirms the system’s sensitivity to structural failure. The accompanying precision of 94.1% indicates a low false-positive rate, avoiding “alarm fatigue” in operational control rooms. In this context, the more diagnostically meaningful metric is Recall, which quantifies the system’s ability to minimize missed alerts. However, as Section 4.5 demonstrates, high statistical recall on hold-out data does not necessarily translate into a predictive lead-time advantage during rapid transients.

4.3 Feature Importance: Validating the Physics Integration

The core hypothesis of this research is that machine learning models for power systems must rely on structural physical metrics rather than memorizing stochastic sensor noise. The global feature importance of the LightGBM model validates this approach.

4.3.1 Structural Dominance




Figure 3: Global feature importance for the LightGBM inference engine, highlighting the dominance of physics-based momentum indicators over raw environmental context.

The global feature importance analysis reveals a critical insight into the model’s decision-making hierarchy. As illustrated in Figure 3, short-term frequency derivatives (rocof_1s at 15.4%) and the absolute grid frequency (9.4%) overwhelmingly dominate the decision splits. In contrast, contextual features such as the OpSDA Wind Ramp Rate and Renewable Penetration exhibit near-zero split-count importance. This finding suggests that for sub-second transient instability prediction, the model prioritizes high-fidelity temporal momentum signals over slower-moving environmental or energy-mix data. This prioritization validates the system’s design as a transient protection mechanism, where the primacy of immediate physical derivatives outweighs the predictive utility of macro-scale meteorological trends.

4.4 Probabilistic Forecasting Reliability

The operational utility of GridGuardian depends on whether its stated risk probabilities align with empirical observation.

4.4.1 Isotonic Recalibration Performance

The recalibrated PICP result of 73.5% still represents a failure of the 80% threshold even after recalibration.




Figure 4: Calibration reliability diagrams demonstrating the alignment of forecasted probability quantiles with empirical observation frequencies.

The calibration analysis (Figure 4) reveals that prior to post-hoc adjustment, the raw model exhibited a significant bias. The application of Isotonic Regression attempts to remap these probabilities; however, the persistent 73.5% PICP failure suggests that recalibration alone cannot compensate for the model’s structural inability to capture the most extreme tail-risk events. While the bounds are “trustworthy” for 73% of observations, they do not yet satisfy the operational requirements for automated EFR intervention without a wider safety margin.

4.5 Forensic AI (XAI) and The August Blackout Reconstruction

The definitive evaluation of the system is its response to the August 9, 2019 blackout. Rather than a “predictive warning,” the system’s performance is better characterized as a **high-fidelity forensic reconstruction** of the collapse.

4.5.1 The Fragility Fingerprint: Reactive Identification




Figure 5: Raw frequency time-series of the August 2019 blackout event overlaid with the LightGBM predicted quantile boundaries, illustrating the reactive alert identification and quantile bound behaviour during the collapse.

Signal Feature	12:00:00 (Stable Noon)	15:52:36 (Pre-Collapse)	15:53:49 (Nadir)	Physical Interpretation
RoCoF Severity Alert	OFF	ON	ON	Detects initial sudden momentum drop
RoCoF Acceleration Alert	OFF	ON	ON	Detects violent non-linear collapse
10s Volatility Alert	OFF	OFF	ON	Detects severe post-fault oscillations
Low Inertia Vulnerability	OFF	ON	ON	Flags insufficient grid momentum
LGBM Quantile Risk ($> 35\%$)	OFF	OFF	ON	Lower quantile bound statistically breached

Table 6: Signal state analysis tracing the multi-physics ‘Fragility Fingerprint’ during the August 9, 2019 blackout.

As documented in Table 6, the ‘Fragility Fingerprint’ was successfully activated during the fault window. However, the primary alert trigger occurred at **15:54:04 UTC**. Given that the

frequency nadir (the point of maximum instability) occurred at **15:53:49 UTC**, this trigger represents a **reactive identification** of the failure rather than a predictive warning. The system provided a 15-second “post-nadir” forensic confirmation. While this is operationally useful for automated incident reporting, it fails to meet the 1.0-second predictive lead-time required for EFR battery deployment during this specific severe event.

4.5.2 SHAP Waterfall Decomposition




Figure 6: SHAP Waterfall decomposition for the reactive alert at 15:54:04 UTC. The model identifies extreme RoCoF as the primary driver, confirming physical consistency but highlighting the lack of predictive “early-warning” indicators.

The SHAP waterfall (Figure 6) confirms that at 15:54:04, the model’s logic was dominated by immediate 1s RoCoF and frequency deviation. The lack of predictive success is likely attributable to the model’s reliance on these high-speed derivatives, which only manifest **after** the initial physical disconnect has occurred. This suggests that for future iterations, the model requires features capable of detecting the pre-fault loss of synchronous momentum prior to the visible frequency drop. This level of granular explainability satisfies the transparency requirements for engineering audits, even when the system fails its primary predictive objective.

Chapter 5: Evaluation and Discussion

5.1 Introduction

This chapter provides a critical analysis of the empirical findings presented in Chapter 4. The objective is to evaluate the operational significance of the GridGuardian system, assess its generalization capabilities across seasonal shifts, and formally justify the necessity of its hybrid architecture. Finally, this chapter critically examines the system’s limitations and proposes concrete avenues for future industrial deployment.

5.2 Methodological Defense: Target Variable and Predictive Potential

A concern for data-driven grid monitoring is the potential for tautology, where the model learns to approximate the physical rules it is meant to predict. The risk of tautological learning is partially mitigated by the $t + 10$ forecasting horizon: the model must project future frequency trajectories, not simply re-evaluate instantaneous conditions. Nevertheless, the reliance on `rocof_1s` in the SHAP analysis (Section 4.3) suggests that the model partially reduces to a reactive derivative detector under transient conditions, which is a limitation documented in Section 5.8. While GridGuardian is trained to forecast a $t + 10$ state based on physical thresholds, its evaluation is grounded in its ability to convert reactive metrics into predictive temporal buffers. Although the prototype fired reactively during the August 2019 blackout (at 15:54:04 UTC), the multi-signal fusion demonstrated a capability to integrate diverse physical signals—such as inertia proxies and RoCoF—into a single decision engine. The system’s contribution is therefore suggested by the implementation evidence as a candidate engineering framework for sub-second data fusion, pending the calibration improvements required before operational deployment.

5.3 Impact Analysis: Operational Viability and Visualization

An operational finding of this research is the system’s ability to reconstruct the precursors of instability, although the predictive margin was limited by the volatility of the August 9 event. To contextualize this result, it must be evaluated against the mechanical deployment constraints of the National Energy System Operator’s (NESO) stability services. Modern Enhanced Frequency Response (EFR) battery systems are mandated to achieve 100% active power injection within 1.0 seconds of an automated trigger.




Figure 7: Hypothetical simulated intervention timeline illustrating the lead-time margin that would be required to trigger EFR battery deployment within the 1.0-second constraint. This figure contextualises the reactive alert produced by the prototype during the August 2019 event against the operational requirement, rather than presenting a measured predictive outcome.

The forensic reconstruction of the August 9 event reveals that the prototype prioritized logical consistency and false-positive suppression over early-warning sensitivity. While the current model iteration achieved high sensitivity in hold-out testing, the reactive nature of the final alert trigger (occurring post-nadir) indicates that the 1.0-second predictive margin required for EFR battery deployment remains an unmet industrial target for this architecture.

To bridge the gap between algorithmic output and operator decision-making, a proof-of-concept Streamlit dashboard was developed as the final visualization layer of the GridGuardian system. This interface presents real-time quantile prediction bands and a dynamic SHAP feature attribution panel. While this dashboard effectively demonstrates the integration of multi-physics signals, it was not subjected to formal usability testing. The author acknowledges that a Human-in-the-Loop evaluation with National Energy System Operator (NESO) controllers would be a prerequisite for actual industrial deployment to ensure the interface facilitates emergency grid response.

5.4 Socio-Economic Impacts and Ethical Considerations

The technical validation of GridGuardian must be contextualized within the landscape of national infrastructure resilience. Power grid instability is not merely a mathematical anomaly; it is a catalyst for societal disruption.

5.4.1 Protection of Infrastructure and Forensic Diagnostics

Grid instability affects vulnerable populations and critical infrastructure. For healthcare providers and electrified rail networks, the 2019 blackout demonstrated sensitivity to frequency nadirs. By providing an interpretable account of the ‘Fragility Fingerprint,’ the GridGuardian system offers grid operators a forensic tool to understand the cascading failure

modes that lead to regional blackouts. While the current prototype did not achieve a predictive lead-time capable of preempting the 2019 event, its ability to provide instantaneous, SHAP-verified diagnostics could accelerate the post-fault recovery of critical infrastructure by isolating the primary physical drivers of the collapse.

5.4.2 Ethical Deployment of Automated Intervention

The transition from predictive alerting to automated intervention introduces significant ethical and legal questions. The GridGuardian architecture addresses this through its “Human-in-the-Loop” explainability. By employing SHAP waterfall decompositions, the system provides grid dispatchers with the physical rationale for every high-risk alert. This transparency is necessary for any AI system deployed in safety-critical national infrastructure, ensuring that automated decisions are accountable, auditable, and physically defensible.

5.5 Out-of-Season Performance and Calibration Drift

A common vulnerability in machine learning models trained on time-series data is seasonal overfitting. The LightGBM model—trained exclusively on August data—was evaluated against a held-out winter testing set (December 2019).

Evaluation Metric	August 2019 (Summer)	December 2019 (Winter)
Mean Absolute Error (Lower)	0.0243 Hz	0.0158 Hz
Pinball Loss ($\alpha = 0.1$) [†]	0.0012	0.0016
Prediction Interval Coverage	73.5%	78.5%

Table 7: Comparative performance and distributional shift evaluation when the summer-trained model is applied to winter grid conditions. [†]Post-recalibration Pinball Loss. The pre-recalibration Pinball Loss baseline (0.0032) is reported in Table 4.

The evaluation reveals that while the winter MAE performance appears superior (0.0158 Hz), the summer-to-winter transfer demonstrates that the PICP remains below the 80% threshold in both seasons. The apparent improvement in winter coverage (78.5% vs. 73.5% in summer) is attributable to structurally lower frequency volatility in December—where demand profiles are more predictable and large renewable ramps are less frequent—rather than genuine generalization gains. A more accurate indicator of performance deterioration is found in the Pinball Loss ($\alpha = 0.1$), which rose from 0.0012 to 0.0016, confirming that the model’s uncertainty quantification deteriorates when the feature distribution shifts away from the summer training regime. This result constitutes a partial falsification of H_1 ’s generalization criterion, confirming that the model cannot be run statically across seasons without periodically recalibrating its quantile bounds against recent historical data to maintain safety margins.

5.6 Global Interpretability and XAI Stability

GridGuardian addresses the “black box” nature of AI through the evaluation of TreeSHAP.

5.6.1 Feature Magnitude and Decision Boundaries




Figure 8: Global SHAP beeswarm plot detailing how feature magnitudes (e.g., extreme negative RoCoF) mathematically push the model towards an unstable frequency prediction.

The analysis confirms that the model correctly interprets the directional physics of the grid. High feature values for RoCoF—representing severe downward acceleration—consistently generate negative SHAP values, dragging the predicted lower bound frequency towards the unstable 49.80 Hz threshold.

5.7 Architectural Justification: Ablation Study

To empirically justify the complexity of the GridGuardian system, an ablation study evaluates the core architectural components in isolation.

Architecture	Inference Latency	Recall	Primary Limitation
Pure LightGBM	0.15s	89%	Overconfidence during unprecedented structural trips.
Pure LSTM	1.45s	95%	Inference latency strictly violates the 1.0s EFR constraint.
Hybrid (Grid-Guardian)	0.20s	99%	Computationally viable; requires reliability scaling.

Table 8: Ablation study comparing the inference latency and predictive recall of isolated architectures vs. the hybrid system.

The ablation study identifies the trade-offs between predictive accuracy and operational latency. While the LSTM baseline demonstrated marginally superior recall, its 1.45-second inference latency renders it non-viable for real-time EFR dispatch. In contrast, the GridGuardian

hybrid architecture achieves a 0.20s inference speed. This performance identifies a design tension: the model is ‘fast enough’ for the grid but currently ‘not reliable enough’ (73.5% PICP) for automated safety interventions. This finding quantifies the performance gap that exists between high-latency deep learning and low-latency ensemble methods.

5.8 System Limitations

Despite its implementation depth, the GridGuardian architecture possesses inherent limitations:

1. **API Latency Dependency:** The current prototype relies on public APIs. A safety-critical system cannot rely on internet-based HTTP requests for real-time control.
2. **Reliability vs. Latency Trade-off:** While the system is fast, the persistent 73.5% PICP failure indicates that reliability was sacrificed for computational speed.
3. **Reactive Blackout Alert:** The forensic reconstruction confirms that the model fired 15 seconds after the nadir, identifying the failure rather than predicting it.
4. **Filter-Induced Lag:** The 5-second smoothing filter used for RoCoF calculation introduces inherent temporal lag, further consuming the predictive window.
5. **Extreme Data Scarcity:** Generalizing across diverse failure modes requires training on multi-year datasets, which were not available for this research.
6. **Lack of Human Validation:** The dashboard and alerting interface were not subjected to formal usability testing with grid operators.
7. **Calibration Leakage Magnitude:** As noted in Chapter 3, the inclusion of August 7-8 in the calibration set introduced a risk of temporal leakage. Forensic estimates suggest this may have artificially inflated the baseline PICP by approximately 1-2%, indicating that true operational reliability on entirely unseen grid topologies may be lower than the reported 73.5%.

5.9 Recommendations for Future Research

To transition GridGuardian into a deployable industrial asset, future research should focus on:

1. **Edge Computing Deployment:** Compiling the engine into C++ for deployment on FPGA hardware located at grid substations to process telemetry directly.
2. **Dynamic Filtering:** Replacing static moving averages with Kalman filtering to estimate grid state with lower inherent lag.
3. **Bayesian Optimization:** Implementing automated frameworks to search the hyperparameter space for more reliable (higher PICP) tree configurations.
4. **High-Frequency Data Integration:** Moving beyond 1Hz telemetry to incorporate 50Hz PMU data for more granular transient detection.

Chapter 6: Conclusion

6.1 Summary of Technical Contributions and Hypotheses Resolution

This research concludes that while a physics-informed, tree-based machine learning architecture offers a computationally efficient and interpretable engine for grid monitoring, it failed to meet the rigorous reliability and lead-time benchmarks required for autonomous grid protection. Despite the implementation of a high-speed inference pipeline (< 0.20 s) and the successful extraction of domain-specific features (RoCoF, OpSDA), the prototype achieved a Prediction Interval Coverage Probability (PICP) of only 73.5% on the primary August 2019 test set (Table 4)—falling short of the 80% operational safety target.

Forensic reconstruction of the August 9, 2019 blackout event revealed that the system functioned as a reactive identifier rather than a predictive early-warning asset, with the primary alert trigger occurring 15 seconds post-nadir (15:54:04 UTC). The available evidence is therefore insufficient to support H_1 : the prototype failed to meet both the 80% PICP criterion and the 1.0-second lead-time criterion during the severe blackout event available for evaluation. H_1 therefore cannot be accepted on the basis of the evidence presented.

The Null Hypothesis (H_0) is resolved in two distinct dimensions. First, H_0 is rejected with respect to the latency and recall comparison: the hybrid LightGBM architecture significantly outperforms the LSTM baseline, achieving a sub-second inference latency of < 0.20 seconds and a recall rate of 99.2% (Table 5), confirming that physics-informed feature engineering yields a structural advantage for time-critical alerting. Second, H_0 is not rejected with respect to industrial safety reliability: the observed PICP of 73.5% (Table 4)—falling below the 80% target—and the confirmed reactive rather than predictive behavior during the August 2019 nadir validate H_0 's claim that the current pipeline cannot consistently meet safety targets across diverse grid conditions. This dual resolution demonstrates that the GridGuardian prototype represents a meaningful engineering advance over purely reactive baselines, while simultaneously establishing the calibration and generalization benchmarks that must be addressed before operational deployment.

The hybrid architecture successfully achieved a 99.2% recall rate on the hold-out test set while maintaining a 0.20-second inference latency, significantly outperforming the deeper LSTM baseline (1.45s). However, the coverage failure indicates that the model's uncertainty estimates were overconfident during the transient blackout window. Global SHAP analysis confirmed that the model correctly prioritized short-term frequency derivatives (rocof_1s) over raw environmental noise, validating the physics-informed design even where the strict statistical reliability thresholds were not met. The results suggest that the “decision window” for automated grid response remains a formidable engineering challenge that requires further reconciliation between algorithmic speed and statistical coverage.

6.2 Professional Reflection

Developing GridGuardian was a transition from textbook theory to the reality of engineering a safety-critical pipeline. One of the biggest technical hurdles was managing the high-resolution frequency data. I initially tried using Pandas, but the memory overhead and performance issues during the temporal alignment of 1-second telemetry data caused constant crashes. This pushed me to learn and implement a high-performance Polars pipeline. The switch to Polars and the use of its `join_asof` function not only solved the performance bottleneck but taught me the value of choosing the right tool for large-scale data engineering.

Beyond the code, the most humbling part of the project was admitting that my model failed to meet its 80% reliability target. In university coursework, you often focus on achieving the highest accuracy, but this project showed me that in the real world—especially in national infrastructure—admitting a model’s limitations is as important as achieving a headline metric. This experience has shifted my perspective from being a student to an engineer who understands that transparency and reliability are critical for safety-critical systems.

Ultimately, the development and evaluation of GridGuardian identify the ‘performance ceiling’ of current ensemble-based grid AI when constrained to 1Hz telemetry. This research contributes a fully implemented software framework, a causal data-merging pipeline, and a forensic auditing methodology to the field of power system security. By documenting the exact failure modes and calibration drift of the prototype, this work provides the empirical foundation required to bridge the gap between theoretical machine learning and the safety requirements of national infrastructure. The project serves as a case study in the practical constraints of physics-informed design, demonstrating that while algorithms can be engineered for speed, operational integrity cannot be compromised by coverage failures — a tension that must be resolved before ensemble methods can be responsibly deployed.

Bibliography

- AEMO. (2017). *Black System in South Australia 28 September 2016* [Technical report]. https://aemo.com.au/-/media/files/electricity/nem/market_notices_and_events/power_system_incident_reports/2017/integrated-final-report-sa-black-system-28-september-2016.pdf
- Bristol, E. H. (1990). Swinging door trending: Adaptive trend recording. *ISA National Conference Proceedings*, 749–753.
- Christ, M., Braun, N., Neuffer, J., & Kempa-Liehr, A. W. (2018). Time series feature extraction on basis of scalable hypothesis tests (tsfresh—A python package). *Neurocomputing*, 307, 72–77.
- Covert, I., Lundberg, S., & Lee, S.-I. (2021). Explaining by removing: A unified framework for model explanation. *Journal of Machine Learning Research*, 22(209), 1–90.
- Gal, Y., & Ghahramani, Z. (2016). Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning. *Proceedings of the 33rd International Conference on Machine Learning (ICML)*, 48, 1050–1059.
- Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.
- Hastie, T., Tibshirani, R., & Friedman, J. (2009). *The elements of statistical learning: data mining, inference, and prediction*. Springer Science & Business Media.
- Homan, S., MacCormac, J., Liaros, C. P. N., & Taylor, P. (2020). The 9 August 2019 power outage in the UK: a review of the event and its implications. *Proceedings of the Institution of Civil Engineers-Energy*, 173(4), 155–162.
- Karniadakis, G. E., Kevrekidis, I. G., Lu, L., Perdikaris, P., Wang, S., & Yang, L. (2021). Physics-informed machine learning. *Nature Reviews Physics*, 3(6), 422–440.
- Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., & Tie-Yan, L. (2017). Lightgbm: A highly efficient gradient boosting decision tree. *Advances in Neural Information Processing Systems*, 3146–3154.
- Koenker, R., & Bassett Jr, G. (1978). Regression quantiles. *Econometrica: Journal of the Econometric Society*, 33–50.
- Kuleshov, V., Fenner, N., & Ermon, S. (2018). Accurate uncertainties for deep learning using calibrated regression. *International Conference on Machine Learning*, 2796–2804.
- Kundur, P., Balu, N. J., & Lauby, M. G. (1994). *Power system stability and control*. McGraw-Hill.
- Liao, W., Bak-Jensen, B., Pillai, J. R., Wang, Y., & Wang, Y. (2021). A review of graph neural networks and their applications in power systems. *Journal of Modern Power Systems and Clean Energy*, 10(2), 345–360.

- Lim, B., Arik, S. O., Loeff, N., & Pfister, T. (2021). Temporal fusion transformers for interpretable multi-horizon time series forecasting. *International Journal of Forecasting*, 37(4), 1748–1764.
- Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, 4765–4774.
- Machlev, R., Heistrene, L., Perl, M., Levy, K. Y., Belikov, J., & Levron, Y. (2022). Explainable Artificial Intelligence (XAI) techniques for energy and power systems: Review, challenges and opportunities. *Energy and AI*, 9, 100169.
- Milano, F., Dörfler, F., Hug, G., Hill, D. J., & Verbič, G. (2018). Foundations and challenges of low-inertia systems. *2018 Power Systems Computation Conference (PSCC)*, 1–25.
- National Grid ESO. (2020). *Operability Strategy Report 2020* [Technical report]. <https://www.nationalgrideso.com/document/183031/download>
- Ofgem and BEIS. (2019). *Report on the Events of 9 August 2019 and System Operator Actions* [Technical report]. <https://www.ofgem.gov.uk/publications/investigation-9-august-2019-power-outage>
- Qiu, X., Gu, W., Wu, J., & Wang, P. (2020). Ensemble machine learning for power system frequency response prediction. *IEEE Transactions on Industrial Informatics*, 17(4), 2556–2565.
- Raissi, M., Perdikaris, P., & Karniadakis, G. E. (2019). Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378, 686–707.
- Shapley, L. S. (1953). A value for n-person games. *Contributions to the Theory of Games*, 2(28), 307–317.
- Slack, D., Hilgard, S., Jia, E., Singh, S., & Lakkaraju, H. (2020). Fooling lime and shap: Adversarial attacks on post hoc explanation methods. *Proceedings of the AAAI/ACM Conference on AI, Ethics, And Society*, 180–186.
- Tielens, P., & Van Hertem, D. (2016). The relevance of inertia in power systems. *Renewable and Sustainable Energy Reviews*, 55, 999–1009.
- Vovk, V., Gammerman, A., & Shafer, G. (2005). *Algorithmic learning in a random world*. Springer Science & Business Media.
- Zhao, J., Gomez-Exposito, A., Netto, M., Mili, L., Abur, A., Terzija, V., Kamwa, I., Pal, B., Singh, A. K., Qi, J., Huang, Z., & Meliopoulos, A. P. S. (2019). Power System Dynamic State Estimation: Motivations, Definitions, Methodologies, and Future Work. *IEEE Transactions on Power Systems*, 34(4), 3029–3043.
- Zhou, H., Zhang, S., Peng, J., Zhang, S., Li, J., Xiong, H., & Zhang, W. (2021). Informer: Beyond efficient transformer for long sequence time-series forecasting. *Proceedings of the AAAI Conference on Artificial Intelligence*, 35(12), 11106–11115.

Chapter 7: Appendices

7.1 Appendix A: Initial Project Proposal

Initial Project Specification (BSc 2025-26)

Student Number: 2604383

Proposed Title: Forecasting UK Frequency Grid Instability: A Data-Driven Alert System

Proposed Aim: Predict UK grid instability from weather and frequency data for proactive renewable energy alert system.

Proposed Objectives:

By the end of this project, I will be able:

1. To research and select a robust proxy variable (e.g., System Frequency deviation $> \pm 0.2$ Hz) for grid instability using National Grid ESO data.
2. To build a demo of data pipeline to integrate, clean, and feature-engineer National Grid ESO and HadUK-Grid datasets using polars.
3. To benchmark three forecasting models (SARIMAX, LightGBM, LSTM) using RMSE, MAE, and training time.
4. To incorporate probabilistic forecasting (quantile regression) to support uncertainty-aware alerts.
5. To develop a proof-of-concept decision-support system integrating a rule-based alert engine and Streamlit dashboard.
6. To validate the alert system on historical data, including the August 9, 2019 blackout event.
7. Critically evaluate the system's performance using SHAP values and assess deployment viability.

Proposed Rationale:

The UK's transition to renewable energy introduces significant grid volatility. Currently, grid operators use costly, high-carbon reactive measures to prevent blackouts. This project addresses this gap by developing a proactive solution that forecasts instability before it happens, allowing for planned, cost-effective, and low-carbon interventions.

Facilities Required:

Python (Polars, Pandas, Scikit-learn, LightGBM, TensorFlow), Streamlit, HadUK-Grid, National Grid ESO API.

Project Reflection & Outcomes:

The project evolved from the initial proposal's focus on SARIMAX and LightGBM point forecasting to a **Physics-Informed LightGBM Quantile Regression** model, with an LSTM Monte Carlo Dropout model serving as a probabilistic baseline. SARIMAX was found to be insufficient for capturing the non-linear, high-frequency transients observed during grid disturbances. The final implementation incorporates physics-informed features (Swing Equation logic, OpSDA wind compression) and a multi-signal alert fusion engine to provide a robust dashboard solution.

Achieved Technical Deliverables:

Original Objective	Final Implementation Outcome
1. Select robust proxy	Established instability threshold at 49.85 Hz with physics-informed triggers (RoCoF & Acceleration).
2. Data pipeline	Engineered a high-performance Polars ETL pipeline merging 1s frequency, HadUK weather, and inertia data.
3. Benchmark models	LightGBM and LSTM benchmarked; Result: LightGBM quantile regression achieved sub-second inference, outperforming the deeper LSTM baseline.
4. Probabilistic alerts	Implemented Quantile Regression ($\alpha=0.1$ to 0.9) with post-hoc Isotonic Calibration for reliable uncertainty.
5. Streamlit Dashboard	Developed real-time "Control Room" interface with probabilistic bounds and SHAP-based risk driver analysis.
6. Blackout Validation	Validated against August 9, 2019 event; provided high-fidelity reactive identification and post-nadir forensic confirmation.
7. Critical Evaluation	Conducted SHAP global/local importance analysis and out-of-season robustness testing on Winter data.

7.2 Appendix B: Ethical Approval

This project utilizes publicly available open datasets from National Grid ESO and Open-Meteo, classifying it as primary quantitative analysis. As the study does not involve human subjects or the collection of sensitive personal data, it qualifies for the internal ethical approval process for non-human research.

The following documentation confirms the internal ethical clearance obtained for the study:

11/26/25, 11:55 AM

CN6000 Internal Ethical Approval Process 2025

ETHICS CN6000 Internal Ethical Approval Process 2025

Only complete this form if you have already received approval for the Proposal. If you complete this form before getting the final approval of your proposal, this form will be deleted, and you will be required to do it again!

Do you need internal approval?

This section aims to check if you need internal approval in the first place. If you are not applying field research, then you may not need internal or external approval.

1. Have you received the final Proposal Approval? That is the proposal has been reviewed by the 'second supervisor' and it has been approved? *

☒ Yes

☐ No

☐ Maybe

2. My study will not involve collecting any data from humans. It is either just based on published datasets that I did not create or there are no datasets at all part of my study. *

☒ True

☐ False

<https://forms.office.com/Pages/ResponseDetailPage.aspx?d=aaKUX5aud0etRi0Z2BQW1qS2Lw3259NmgdTAKKnKIJUMU8yVThTWVRaV0ZQWV...>

Figure 9: Ethical Approval Form (Page 1)

11/26/25, 11:55 AM CN6000 Internal Ethical Approval Process 2025

3. Please state your topic title: * 

GridGuardian: Evaluating Physics-Informed Machine Learning for Grid Instability Early-Warning Systems

4. Please state your supervisor's name (Spell it correctly). * 

Ms. Dhara Parekh

5. If you answered **Yes** to question 1 and **True** to Question 2 - You are not required to proceed. Just click NEXT then Submit without the need to answer the remaining questions. Otherwise, please proceed to section 2. * 

☒ True

☐ False

This section aims to see if we can approve your proposal internally. 

Please read the questions carefully. Feel free to discuss these questions with your supervisor. If you find out that you have to apply for approval by the Research Ethics Committee, please discuss with your supervisor if there is a way to modify your dissertation in a way that does not take away from the value of your work.

6. I intend to plan and conduct my study as a pilot test of the research only. 

☐ Yes

☐ No

<https://forms.office.com/Pages/ResponseDetailPage.aspx?id=aaKLu5sudDetRi022BGW1qS2Ltw3259NmgdTakKknKJUMU8yVThTWV/RaV0ZOWV...> 2/5

Figure 10: Ethical Approval Form (Page 2)

11/26/25, 11:55 AM CN6000 Internal Ethical Approval Process 2025

☐ Maybe

7. I intend to conduct my field research as pilot test on UEL students or staff only. ☐

☐ Yes

☐ No

☐ Maybe

8. I will debrief participants before the study and mention that their contribution is entirely voluntary. ☐

☐ Yes

☐ No

☐ Maybe

9. I will not be collecting any personal data such as names, addresses, telephone numbers, emails...etc. ☐

☐ Yes

☐ No

☐ Maybe

10. I will not be asking any sensitive questions that can cause distress - such as mental health related or emotional questions. ☐

☐ Yes

<https://forms.office.com/Pages/ResponseDetailPage.aspx?id=aakUv5aud0eIRi0Z2BGW1qS2Lw3259NmgdTaKKknKJUMUbyVTnTWWBaV0ZQWV...>

Figure 11: Ethical Approval Form (Page 3)

11/26/25, 11:55 AM CN6000 Internal Ethical Approval Process 2025

☐ No

☐ Maybe

11. I will be getting approval of the format and content of my questions from my supervisor before I proceed with my field research. ☐

☐ Yes

☐ No

☐ Maybe

12. I will have a 'consent' form for participants to be part of their study and be informed of their rights. ☐

☐ Yes

☐ No

☐ Maybe

13. I will store the data I collect on UEL OneDrive. No other copy of it will be stored elsewhere. ☐

☐ Yes

☐ No

☐ Maybe

14. I will not use the data for any other intended purpose and delete it at the end of the term. ☐

<https://forms.office.com/Pages/ResponseDetailPage.aspx?id=aakUv5aud0eiRi0Z2BGW1qS2Lw3259NmgdTaKKknKJUMUbyVTnTWWBaV0ZQWV...>

Figure 12: Ethical Approval Form (Page 4)

11/26/25, 11:55 AM CN6000 Internal Ethical Approval Process 2025

☐ Yes

☐ No

☐ Maybe

15. If you have answered No or Maybe to any of the questions above, you cannot proceed with internal Ethical approval. I would advise you to review those particular questions with your supervisor before proceeding. If you both agree that you want to keep these answers, then please indicate below that you are happy to proceed. 

☐ All my answers are Yes. I have completed the internal approval and promise to abide by it.

☐ Not all my answers are yes. I have spoken to my supervisor and we both agreed that I should apply for form B Ethical Approval process.

16. I understand that violating any of the terms listed above could result in an automatic fail and possibly academic misconduct proceedings. 

☐ I confirm

 Microsoft 365

This content is created by the owner of the form. The data you submit will be sent to the form owner. Microsoft is not responsible for the privacy or security practices of its customers, including those of this form owner. Never give out your password.

Microsoft Forms | AI-Powered surveys, quizzes and polls [Create my own form](#)

[Privacy and cookies](#) | [Terms of use](#)

<https://forms.office.com/Pages/ResponseDetailPage.aspx?id=aaKUx5sud0etRi0Z2BGW1qS2Ljw3259Nm9dTakKknKJUMU8yVThTWVRaV0ZQWV...> 5/5

Figure 13: Ethical Approval Form (Page 5)

7.3 Appendix C: Dashboard Outputs

This appendix provides a visual walkthrough of the GridGuardian “Control Room” dashboard. The interface translates complex, high-resolution machine learning outputs—such as LightGBM quantile regression bounds, multi-signal physics heuristics, and SHAP-based feature importance—into a user-friendly, real-time diagnostic tool for grid operators.

The Control Room dashboard is structured to provide intuitive, real-time oversight of grid stability. Key features include:

- **Time Navigation (Sidebar):** Allows operators to jump to specific dates, scrub through time with 1s/5s precision, or use autoplay to seamlessly replay historical transients like the August 2019 blackout.
- **Top KPI Panel:** Provides an immediate summary of the system state, displaying real-time frequency, predicted bounds, and an aggregated status banner (e.g., Stable, Uncertainty, Alert) alongside the Time to Alert.
- **Real-Time Monitor (Main Chart):** Visualizes the raw telemetry against the LightGBM predicted quantile bounds (uncertainty ribbon), immediately highlighting when the grid is approaching physical thresholds.
- **Risk Drivers (XAI Chart):** Translates complex model features into an interpretable SHAP waterfall chart, instantly isolating whether inertia, wind ramps, or RoCoF is driving the current instability risk.

7.3.1 C.1: Normal Operational State (Pre-Fault)




Figure 14: Dashboard displaying stable grid conditions prior to the fault event.

System State: Stable (Frequency near 50.0 Hz).

Insight: The uncertainty bounds are narrow and the risk drivers show minimal physical stress, indicating healthy grid inertia and nominal volatility.

7.3.2 C.2: High Volatility & Model Uncertainty




Figure 15: The system detecting early-warning turbulence, triggering a “High Model Uncertainty” state.

System State: High Model Uncertainty (Pre-Alert Buffer).

Insight: As turbulence increases, the LightGBM confidence intervals dynamically widen. This alerts operators to elevated volatility before the absolute frequency breaches critical safety limits.

7.3.3 C.3: Critical Instability Alert (August 9, 2019 Nadir)




Figure 16: Dashboard alert during the August 9, 2019 event. Note: This visualization represents a high-sensitivity simulation mode (captured at 15:52:34 UTC) used for demonstration; the primary operational alert reported in Chapter 4 occurred at 15:54:04 UTC.

System State: Instability Alert (Simulation Replay).

Insight: The multi-signal fusion engine converges as frequency drops and RoCoF plummets. While the operational model prioritized false-positive suppression (leading to a reactive alert), this simulation demonstrates the engine’s potential for earlier detection if sensitivity thresholds are relaxed.

7.3.4 C.4: Post-Fault Recovery Monitoring




Figure 17: The recovery phase, demonstrating the model tracking the post-nadir stabilization.

System State: Post-Fault Recovery.

Insight: The uncertainty bounds expand to capture massive variance during the unprecedented transient phase, while the LSTM probability timeline (bottom chart) tracks the grid’s gradual stabilization.

7.3.5 C.5: Model Health & Calibration Diagnostics




Figure 18: The “Model Health” tab providing continuous transparency into the ML engine’s performance.

System State: Diagnostic Oversight.

Insight: Operators can monitor rigorous real-time statistics (Pinball Loss, PICP Coverage, Confusion Matrix). This ensures the ML engine remains a transparent, trusted advisory tool rather than an opaque oracle.

7.4 Appendix D: Code Snippets

This appendix presents the core implementation logic for the GridGuardian system, providing a technical reference for the feature engineering, probabilistic modeling, and real-time dashboard engine. Note: These snippets are provided as abbreviated illustrations of the primary logic; for the full production-ready implementation and environment configurations, refer to the project's GitHub repository.

7.4.1 D.1: Physics-Informed Feature Engineering

The following snippets detail the physics-informed feature engineering pipeline, including the Optimised Swinging Door Algorithm (OpSDA) implementation and the derivation of causal multi-window RoCoF and wind ramp rates.

Optimised Swinging Door Algorithm (Implementation: `opsda.py`)

```
def compress(data, width):
    """
    Compresses a list of (timestamp, value) tuples using the Swinging Door
    Algorithm.
    """
    if not data:
        return []

    compressed_data = [data[0]]
    start_point_index = 0

    for i in range(1, len(data)):
        current_point = data[i]
        pivot_point = data[start_point_index]

        # Form a "door" from the pivot point to the current point
        upper_bound_slope = (pivot_point[1] + width - current_point[1]) /
(pivot_point[0] - current_point[0]) if pivot_point[0] != current_point[0] else
float('inf')
        lower_bound_slope = (pivot_point[1] - width - current_point[1]) /
(pivot_point[0] - current_point[0]) if pivot_point[0] != current_point[0] else
float('-inf')

        # Check all intermediate points
        for j in range(start_point_index + 1, i):
            intermediate_point = data[j]
            slope = (pivot_point[1] - intermediate_point[1]) / (pivot_point[0]
- intermediate_point[0]) if pivot_point[0] != intermediate_point[0] else
float('inf')

            if slope > upper_bound_slope or slope < lower_bound_slope:
                compressed_data.append(data[i-1])
                start_point_index = i - 1
```

```
break
```

```
compressed_data.append(data[-1])
return compressed_data
```

Wind Ramp Rate and Causal RoCoF (Selection from: feature_engineering.py)

```
def calculate_wind_ramp_rate(df):
    """Abbreviated illustration of Polars-based OpSDA alignment."""
    weather_data = df.select(["timestamp",
                              "wind_speed"]).unique(subset=["timestamp"])
    weather_data = weather_data.with_columns(
        unix_ts = pl.col("timestamp").dt.timestamp("s")
    )

    data_tuples = list(weather_data.select(['unix_ts',
                                           'wind_speed']).iter_rows())
    compressed = opsd.compress(data_tuples, width=OPSDA_WIDTH)

    compressed_df = pl.DataFrame(compressed, schema=["timestamp",
                                                    "wind_speed"])
    compressed_df = compressed_df.with_columns(
        wind_ramp_rate = pl.col("wind_speed").diff() /
        (pl.col("timestamp").diff().dt.nanoseconds() / 1e9)
    )
    return df.join_asof(compressed_df, on="timestamp", strategy="backward")

def create_features(df):
    """Selection of causal feature engineering logic in Polars."""
    df = df.with_columns([
        pl.col("grid_frequency").diff(1).fill_null(0).alias("rocof_1s"),
        ((pl.col("grid_frequency") - pl.col("grid_frequency").shift(5)) /
        5.0).fill_null(0).alias("rocof_5s"),
        ((pl.col("grid_frequency") - pl.col("grid_frequency").shift(10)) /
        10.0).fill_null(0).alias("rocof_10s"),
    ])

    # Physics-informed risk signals
    df = df.with_columns(
        wind_power_proxy = pl.col("wind_speed").pow(3).mul(3.0).clip(0, 3000),
        rocof_smooth =
    pl.col("rocof_1s").rolling_mean(window_size=5).fill_null(0)
    )

    df = df.with_columns(
        renewable_penetration_ratio = (pl.col("wind_power_proxy") /
    pl.col("demand_proxy").replace(0, 35000)).clip(0, 1)
    )
```

```

df = df.with_columns(
    rocof_inertia_risk = pl.col("rocof_smooth").abs() *
pl.col("renewable_penetration_ratio")
)
return df

```

7.4.2 D.2: Probabilistic Model Logic

This section presents the probabilistic model logic, showcasing the quantile regression setup with weighted sampling to prioritize tail-risk events and the post-hoc isotonic recalibration routine.

Quantile Regression and Weighted Sampling (Selection from: `model_trainer.py`)

```

def train_quantile_model(df, alpha):
    params = LGBM_QUANTILE_PARAMS.copy()
    params['objective'] = 'quantile'
    params['alpha'] = alpha

    model = lgb.LGBMRegressor(**params)

    # Upweight samples where target frequency is low (approaching instability)
    sample_weights = np.where(y_train < 49.95, 10.0,
                              np.where(y_train < 50.00, 3.0,
                                      1.0))
    model.fit(X_train, y_train, sample_weight=sample_weights)

    # Evaluate Pinball Loss on test set
    y_pred_quantile = model.predict(X_test)
    loss = pinball_loss(y_test, y_pred_quantile, alpha)
    print(f"Pinball Loss (alpha={alpha}): {loss:.4f}")
    return model, X_test, y_test

```

LSTM MC Dropout Comparator (Selection from: `model_trainer.py`)

```

def train_lstm_quantile_comparator(df_processed, n_mc_samples=25):
    inputs = tf.keras.Input(shape=(LSTM_TIME_STEPS, train_scaled.shape[1]))
    x = tf.keras.layers.LSTM(50, return_sequences=False)(inputs)
    # Dropout kept ON during inference (training=True) for MC Dropout
    x = tf.keras.layers.Dropout(0.2)(x, training=True)
    outputs = tf.keras.layers.Dense(1)(x) # regression, not classification

    model = tf.keras.Model(inputs, outputs)
    model.compile(optimizer='adam', loss='mae')

    # Stochastic forward passes
    tiled_inputs = tf.repeat(batch[0], n_mc_samples, axis=0)
    tiled_preds = model(tiled_inputs, training=True)

```

```

batch_samples = tf.reshape(tiled_preds, (batch_size, n_mc_samples, 1))
batch_mean = tf.reduce_mean(batch_samples, axis=1)
return model, scaler, batch_mean.numpy()

```

Post-hoc Quantile Recalibration (calibration.py)

```

def fit_calibrator(y_true, y_pred_quantile, alpha):
    """Post-hoc recalibration using isotonic regression."""
    calibrator = IsotonicRegression(out_of_bounds='clip')
    calibrator.fit(y_pred_quantile, y_true)
    return calibrator

def calibrate_predictions(calibrator, y_pred_quantile):
    return calibrator.predict(y_pred_quantile)

```

7.4.3 D.3: Dashboard Engine & Multi-Signal Alert Logic

The dashboard engine logic illustrates the high-performance model monitoring and the multi-physics alert fusion engine that powers the real-time Control Room interface.

Cache Design and Alert Fusion (Selection from: app.py)

```

# Persists heavy model objects across user sessions
@st.cache_resource
def load_models(model_stamp: tuple):
    lower_model = joblib.load(LGBM_QUANTILE_LOWER_PATH)
    upper_model = joblib.load(LGBM_QUANTILE_UPPER_PATH)
    classifier_model = joblib.load(LGBM_MODEL_PATH)
    lstm_model = tf.keras.models.load_model(LSTM_MODEL_PATH)
    return lower_model, upper_model, classifier_model, lstm_model

# Variable Extraction from Current State
rocof_now = current_row.get('rocof_smooth', current_row.get('rocof', 0.0))
rocof_accel = current_row.get('rocof_accel', 0.0)
volatility = current_row.get('volatility_10s', 0.0)
freq_now = current_row['grid_frequency']
ren_pen = current_row.get('renewable_penetration_ratio', 0.0)

# Physics-Informed Multi-Signal Alert Fusion
rocof_alert = (rocof_now < -0.015) and (freq_now < 50.05)
accel_alert = (rocof_accel < -0.005)
volatility_alert = (volatility > 0.02) and (freq_now < 50.1)
renewable_stress = (ren_pen > 0.15) and (rocof_now < -0.01)
freq_boundary = freq_now < 49.95

# Score-based fusion (convergence of signals)
signal_count = sum([
    rocof_alert, accel_alert, volatility_alert,
    renewable_stress, freq_boundary,
    classifier_prob > 0.35, lstm_prob > 0.35,

```

1)

```
emergency_trigger = (signal_count >= 3) or (freq_now < alert_threshold_hz) or
(lower_bound_pred < alert_threshold_hz)
```

```
# Intervention Simulator (Swing Equation Implementation)
#  $\Delta f = (\Delta P \times f_0) / (2 \times H \times S_{base})$ 
if synthetic_inertia_mw > 0:
    swing_delta_f = (synthetic_inertia_mw * NOMINAL_FREQ) / (2 *
SYSTEM_INERTIA_H * TOTAL_SYSTEM_CAPACITY)
    lower_bound_pred = lower_bound_raw + swing_delta_f
```

7.4.4 D.4: Global Configuration

The global configuration file defines the system thresholds, physics constraints, and model hyperparameters used throughout the GridGuardian pipeline.

System Thresholds and Feature Set (config.py)

```
TTA_SECONDS = 10
OPSDA_WIDTH = 0.5
LAG_INTERVALS_SECONDS = [1, 5, 60]

LGBM_FEATURE_COLS = [
    "grid_frequency",
    "rocof_1s", "rocof_5s", "rocof_10s", "rocof_30s", "rocof_accel",
    "rocof_smooth", "rocof",
    "volatility_10s", "volatility_30s", "volatility_60s",
    "wind_speed", "wind_power_proxy", "demand_proxy", "wind_ramp_rate",
    "solar_radiation", "hour", "renewable_penetration_ratio",
    "inertia_value", "inertia_roc", "low_inertia_flag", "rocof_inertia_risk"
] + [f"lag_{lag}s" for lag in LAG_INTERVALS_SECONDS]

LGBM_QUANTILE_PARAMS = {
    'n_estimators': 300,
    'learning_rate': 0.03,
    'max_depth': 6,
    'num_leaves': 31,
    'min_child_samples': 50,
    'subsample': 0.8,
    'colsample_bytree': 0.8,
    'reg_alpha': 0.1,
    'reg_lambda': 1.0,
}
```